

TELEMED AI v4

Master Technical, Machine Learning, Clinical Decision-Support & Software Architecture
Defense Report

Author / Lead: Swaran Gunda | Release Baseline: v4.0-final | Date: August 31, 2026

TELEMED AI v4: MULTIMODAL METABOLIC DECISION SUPPORT PLATFORM

Master Technical, Machine Learning, Clinical Decision-Support & Software Architecture Defense Report

Version: v4.0-final Production Baseline | Status: Release-Gate Verified & Live Deployed

SECTION 1: EXECUTIVE SUMMARY

1.1 Project Title & Core Definition

- **Formal Project Title:** TeleMed AI v4 — Generative AI-Assisted Multimodal Clinical Decision-Support & Telemedicine Platform.
- **One-Line Definition:** An enterprise-grade, privacy-first telemedicine platform combining 7-pathway multimodal machine learning (Clinical, Continuous Glucose/Wearable, and 16S Gut Microbiome biomarkers), Unified TreeSHAP explainability, and evidence-grounded Medical Retrieval-Augmented Generation (RAG) to provide calibrated metabolic risk stratification without cross-modal data imputation.

1.2 Problem Being Solved & Motivation

Metabolic and chronic non-communicable diseases—specifically Type 2 Diabetes, Prediabetes, High Adiposity Risk (Obesity), Metabolic Syndrome, and Non-Alcoholic Fatty Liver Disease (NAFLD)—represent a global epidemic causing over 70% of premature non-communicable deaths worldwide. Traditional telemedicine platforms suffer from critical deficiencies:

1. **Unimodal Isolation:** Most platforms evaluate only basic clinical labs or self-reported questionnaires, ignoring continuous glycemic dynamics and gut microbiome dysbiosis.
2. **Forced Data Imputation / Hallucination:** When patient data is incomplete (e.g. missing wearable or gut microbiome test), existing systems either fail completely or impute synthetic averages, introducing severe data contamination and false clinical certainty.
3. **Black-Box AI:** Opaque neural networks output arbitrary risk numbers without transparent biomarker-level attribution, leading to clinician distrust.
4. **Ungrounded Generative AI:** Consumer LLMs hallucinate medical guidance when answering patient queries without grounding in peer-reviewed clinical guidelines.

1.3 Target Users & Clinical Stakeholders

- **Patients:** Individuals seeking preventive metabolic assessment, multimodal biomarker logging, automated explainable risk insights, appointment scheduling, and encrypted telemedicine consultations.
- **Clinicians & Healthcare Providers:** Physicians, endocrinologists, and gastroenterologists requiring objective risk stratification, TreeSHAP biomarker drivers, AI-synthesized clinical summaries with guideline citations, and an integrated consultation workspace.
- **Health System Administrators:** Clinical compliance officers and compliance auditors monitoring system throughput, doctor credential verification, and cryptographic immutable audit ledgers.

1.4 Core Innovations & Technical Architecture

1. **Exact 7-Pathway Routing Architecture:** Accommodates all $2^3 - 1 = 7$ mathematical permutations of available patient data:
 - * **Unimodal Pathways:** Pathway 1 (\$C\$: Clinical alone), Pathway 2 (\$W\$: Wearable alone), Pathway 3 (\$G\$: Gut Microbiome alone).
 - * **Bimodal Pathways:** Pathway 4 (\$C+W\$: Clinical + Wearable), Pathway 5 (\$C+G\$: Clinical + Gut), Pathway 6 (\$W+G\$: Wearable + Gut).
 - * **Trimodal Pathway:** Pathway 7 (\$C+W+G\$: Clinical + Wearable + Gut Fusion).
2. **Zero-Imputation Missing-Modality Guarantee:** Modalities that are not provided by the patient remain strictly null (NOT PROVIDED), eliminating cross-modal leakage and false feature synthesis.

3. **Exact 5-Target Multimodal Disease Spectrum:** Calibrated predictions for exactly five metabolic conditions: Type2_Diabetes, Prediabetes, High_Adiposity_Risk, Metabolic_Syndrome, and NAFLD.
 4. **Unified TreeSHAP Explainability:** Fast tree-ensemble feature attribution calculating exact local Shapley values ($\mathcal{O}(\text{TLD}^2)$ complexity) for every prediction.
 5. **Deterministic Evidence-Grounded Medical RAG:** Vector-indexed clinical guidelines (ADA, EASD, AHA/ACC, AASLD, AGA) synthesized through FAISS vector retrieval with explicit evidence citations.
 6. **Multi-Portal Zero-Cost Production Stack:** Vercel (React 18 Frontend) + Render (FastAPI ASGI + 7-Pathway ML Engine) + Neon Cloud (Serverless PostgreSQL 17).
-

SECTION 2: PROBLEM STATEMENT, OBJECTIVES & ABSTRACT

2.1 Formal Problem Title

Design, Implementation, and Empirical Validation of a Zero-Imputation 7-Pathway Multimodal Machine Learning and Retrieval-Augmented Generation Architecture for Metabolic Syndrome Risk Stratification in Distributed Telemedicine Platforms.

2.2 Detailed Problem Statement

In modern clinical practice, metabolic disorders develop silently through interconnected pathophysiological axes: systemic insulin resistance (reflected in fasting glucose, HbA1c, and lipid panels), glycemic variability (captured by continuous glucose monitors and autonomic heart rate metrics), and gut mucosal dysbiosis (quantified by 16S rRNA relative bacterial abundances). Despite the known cross-talk between these biological layers, clinical decision support in telemedicine remains fragmented. When patients present with partial biomarker panels, conventional systems either break or fabricate values via mean/median imputation, severely violating medical reliability standards.

2.3 Research & Engineering Objectives

1. **Objective 1 (Multimodal Integrity):** Engineer 7 independent, dynamically routed inference pipelines that process any combination of Clinical, Wearable, and Gut Microbiome data without synthetic feature imputation.
2. **Objective 2 (Calibrated Stratification):** Train and evaluate expert ensembles across exactly 5 metabolic targets, maximizing Macro F1 and ROC-AUC.
3. **Objective 3 (Clinical Explainability):** Implement Unified TreeSHAP attribution delivering top risk-increasing and risk-decreasing biomarkers per patient in ≤ 400 ms.
4. **Objective 4 (Grounded Decision Support):** Implement a Medical RAG pipeline retrieving verified medical guidelines to eliminate generative AI hallucinations.
5. **Objective 5 (Enterprise Software & Security):** Deliver role-based access control (RBAC), multi-tenant IDOR protection, cryptographic audit logging, and automated CI/CD deployment on a zero-cost production cloud.

2.4 Functional & Non-Functional Requirements

• Functional Requirements:

- * Automated PDF/OCR lab report extraction and structured biomarker parsing.
- * 7-pathway dynamic inference with missing modality preservation.
- * TreeSHAP feature attribution waterfall generation.
- * Grounded clinical report generation with guideline citations.
- * Multi-portal workspaces for Patients, Doctors, and Administrators.
- * Real-time consultation booking, messaging, and appointment management.

• Non-Functional Requirements:

- * **Latency:** Single-pathway ML inference ≤ 60 ms; RAG synthesis ≤ 500 ms.
- * **Memory Ceiling:** Total resident set size ≤ 512 MB for zero-cost cloud compatibility.
- * **Security:** Strict JWT authentication, cryptographic audit hashing, and zero cross-tenant IDOR leakage.
- * **Reliability:** 100% test pass rate across unit, integration, and security regression suites.

2.5 Academic Abstract

This report presents **TeleMed AI v4**, an end-to-end clinical decision-support and telemedicine system addressing metabolic disease stratification across Clinical, Wearable Continuous Glucose Monitoring (CGM), and Gut Microbiome modalities. Traditional machine learning models in digital health rely on unimodal inputs or perform uncontrolled feature imputation when facing incomplete multimodal records. TeleMed AI v4 introduces a 7-pathway routing architecture that executes specialized expert models for unimodal (\$C\$, \$W\$, \$G\$), bimodal (\$C+W\$, \$C+G\$, \$W+G\$), and trimodal (\$C+W+G\$) inputs, strictly preserving missing modalities as null to eliminate data leakage. The platform evaluates five metabolic targets: Type 2 Diabetes, Prediabetes, High Adiposity Risk (Obesity), Metabolic Syndrome, and NAFLD. Explainability is delivered via Unified TreeSHAP, computing exact local feature attributions, while a Medical Retrieval-Augmented Generation (RAG) engine grounds AI summaries in peer-reviewed clinical guidelines (ADA, EASD, AHA, AASLD, AGA). Built with React 18, FastAPI, and PostgreSQL 17, TeleMed AI v4 is deployed on a zero-cost cloud architecture (Vercel, Render, Neon) and verified through 147 automated tests, sub-400ms inference latencies, and zero-leakage multi-tenant security isolation.

SECTION 3: COMPLETE SYSTEM ARCHITECTURE

3.1 Architecture Overview & End-to-End Component Flow

```
[User Browser: Patient / Doctor / Admin]
■
▼ (HTTPS / WSS)
[Vercel Edge Network: React 18 + Vite SPA]
■■■ Public Pages (Hero, Features, Research, About, Care)
■■■ Patient Portal (Intake, 7-Pathway AI, Records Vault, Consultations)
■■■ Doctor Portal (Credential Upload, Verification, Clinical Workspace)
■■■ Admin Portal (Doctor Ledger Approval, Audit Logs, System Metrics)
■■■ SPA Routing & /api/* Proxy Rewrites
■
▼ (Secure Reverse Proxy HTTPS)
[Render Cloud: FastAPI ASGI Web Service (app.backend.main:app)]
■■■ Security & Governance Layer
■ ■■■ RateLimitingMiddleware (Sliding window IP throttling)
■ ■■■ SecurityHeadersMiddleware (CSP, X-Frame-Options, HSTS)
■ ■■■ Role-Based Access Control (require_role: PATIENT, DOCTOR, ADMIN)
■ ■■■ Cryptographic Hash-Chained Audit Ledger
■
■■■ Multimodal Intake & Normalization Engine
■ ■■■ PDF Parser & Tesseract OCR Engine
■ ■■■ Schema Validation & Biomarker Bounds Checker (v3_schema_validator)
■ ■■■ Missing Modality Detection & Null Masking
■
■■■ 7-Pathway Dynamic Routing Engine (v3_scientific_router.py)
■ ■■■ Pathway 1 (C) ■■> Clinical Expert (XGBoost / LightGBM)
■ ■■■ Pathway 2 (W) ■■> Wearable CGM Expert (LightGBM / LogReg)
■ ■■■ Pathway 3 (G) ■■> Gut 16S Microbiome Expert (RandomForest / XGBoost)
■ ■■■ Pathway 4 (C+W) ■■> Bimodal Clinical+Wearable Stacking
■ ■■■ Pathway 5 (C+G) ■■> Bimodal Clinical+Gut Stacking
■ ■■■ Pathway 6 (W+G) ■■> Bimodal Wearable+Gut Stacking (L2 LogReg Stacker)
■ ■■■ Pathway 7 (C+W+G) ■■> Trimodal Meta-Learner (Multimodal Stacking)
■
■■■ Explainable AI & Clinical Synthesis Engine
■ ■■■ Unified TreeSHAP Engine (15 fitted TreeExplainers across 5 targets)
■ ■■■ Medical RAG Service (FAISS Vector Index, 20 Chunks, 5 Guidelines)
■
■■■ Database & Storage Layer
■■■ SQLAlchemy 2.0 ORM with Connection Pooling (pool_size=30, max_overflow=20)
■■■ Neon Cloud Serverless PostgreSQL 17 (10 Relational Tables)
```

SECTION 4: DATASETS — COMPLETE INVENTORY & SPECIFICATIONS

4.1 Master Dataset Inventory

Dataset Name	Version	Modality	Source / Status	Total Records	Predictors	Target Labels	Train / Val / Test Split
Clinical_Dataset.csv	V4 (Frozen)	Clinical Labs & Vitals	Synthetic Multi-Cohort	20,000	19 features	5 targets	14,000 / 3,000 / 3,000 (70/15/15%)
Wearable_Dataset.csv	V4 (Frozen)	CGM & Physiological	Synthetic Continuous	20,000	10 features	5 targets	14,000 / 3,000 / 3,000 (70/15/15%)
Gut_Microbiome_Dataset.csv	V4 (Frozen)	16S rRNA Taxa	Synthetic Metagenomic	20,000	10 features	5 targets	14,000 / 3,000 / 3,000 (70/15/15%)

4.2 Feature Schema & Physiological Boundaries

A. Clinical Dataset Feature Schema (19 Predictors)

- age (Years, integer, range 18–85, mean=48.2, std=14.1)
- bmi ($\text{\$ ext{kg/m}^2}$, float, range 16.0–52.0, mean=27.4, std=5.8)
- fasting_glucose ($\text{\$ ext{mg/dL}}$, float, range 65.0–280.0, mean=104.5, std=28.3)
- hba1c (% , float, range 4.5–13.5, mean=5.85, std=1.24)
- systolic_bp ($\text{\$ ext{mmHg}}$, integer, range 85–210, mean=124.8, std=16.5)
- diastolic_bp ($\text{\$ ext{mmHg}}$, integer, range 55–125, mean=79.2, std=10.2)
- total_cholesterol ($\text{\$ ext{mg/dL}}$, float, range 110.0–340.0, mean=198.4, std=38.1)
- ldl_cholesterol ($\text{\$ ext{mg/dL}}$, float, range 45.0–240.0, mean=118.6, std=32.4)
- hdl_cholesterol ($\text{\$ ext{mg/dL}}$, float, range 22.0–95.0, mean=49.1, std=13.8)
- triglycerides ($\text{\$ ext{mg/dL}}$, float, range 40.0–550.0, mean=152.3, std=68.2)
- alt ($\text{\$ ext{U/L}}$, float, range 7.0–145.0, mean=28.4, std=18.6)
- ast ($\text{\$ ext{U/L}}$, float, range 8.0–130.0, mean=26.1, std=15.9)
- ggt ($\text{\$ ext{U/L}}$, float, range 8.0–180.0, mean=32.8, std=22.4)
- creatinine ($\text{\$ ext{mg/dL}}$, float, range 0.4–2.8, mean=0.92, std=0.28)
- uric_acid ($\text{\$ ext{mg/dL}}$, float, range 2.5–11.5, mean=5.6, std=1.5)
- hs_crp ($\text{\$ ext{mg/L}}$, float, range 0.1–18.0, mean=2.15, std=2.4)
- insulin_fasting ($\text{\$ \mu ext{U/mL}}$, float, range 2.0–65.0, mean=11.8, std=8.4)
- homa_ir (Index, float, range 0.4–16.5, calculated: $\text{\$ rac{ ext{glucose} imes ext{insulin}}{405}}$, mean=2.94, std=2.18)
- waist_circumference ($\text{\$ ext{cm}}$, float, range 60.0–145.0, mean=88.5, std=14.2)

B. Wearable CGM Dataset Feature Schema (10 Predictors)


- mean_glucose_24h ($\text{\$ ext{mg/dL}}$, float, range 70.0–240.0, mean=108.2, std=24.1)
- glucose_variability_cv (% , float, range 10.0–48.0, mean=22.4, std=6.8)
- time_in_range_70_180 (% , float, range 30.0–100.0, mean=86.5, std=12.4)
- time_above_range_180 (% , float, range 0.0–65.0, mean=9.8, std=10.2)
- time_below_range_70 (% , float, range 0.0–25.0, mean=3.7, std=3.9)
- resting_heart_rate ($\text{\$ ext{bpm}}$, integer, range 45–115, mean=68.4, std=9.5)
- hrv_sdnn ($\text{\$ ext{ms}}$, float, range 15.0–110.0, mean=48.2, std=16.8)
- sleep_duration_hours (Hours, float, range 3.5–11.0, mean=7.1, std=1.2)
- daily_step_count (Steps, integer, range 800–28000, mean=7450, std=3620)

10. sedentary_hours (Hours, float, range 3.0–16.0, mean=8.8, std=2.3)
- #### C. Gut Microbiome Dataset Feature Schema (10 Relative Abundance Predictors)
1. Bifidobacterium (% , float, range 0.5–28.0, mean=8.4, std=4.6)
 2. Lactobacillus (% , float, range 0.2–18.0, mean=4.2, std=2.8)
 3. Akkermansia (% , float, range 0.0–12.0, mean=2.8, std=2.1)
 4. Faecalibacterium (% , float, range 1.0–24.0, mean=9.1, std=4.2)
 5. Bacteroides (% , float, range 8.0–45.0, mean=24.6, std=7.5)
 6. Firmicutes (% , float, range 15.0–65.0, mean=38.2, std=9.4)
 7. Prevotella (% , float, range 0.5–32.0, mean=7.8, std=6.1)
 8. Roseburia (% , float, range 0.2–14.0, mean=3.9, std=2.4)
 9. Ruminococcus (% , float, range 0.5–16.0, mean=4.8, std=2.9)
 10. Enterobacteriaceae (% , float, range 0.0–15.0, mean=1.8, std=2.2)

4.3 Master Patient Split Protocol

- **Split File:** archive/legacy_datasets/expert_models_splits/patient_split.csv
 - **Partition Size:** Exactly 20,000 unified patient IDs partitioned into:
- * **Training Set:** 14,000 patients (70.0%)
- * **Validation Set:** 3,000 patients (15.0%)
- * **Test Set:** 3,000 patients (15.0%)
- **Leakage Defense:** The exact same patient IDs are indexed across Clinical, Wearable, and Gut datasets. No patient in the test set was ever used in training or hyperparameter tuning.

SECTION 5: DATASET EVOLUTION V1 → V2 → V3 → V4



5.1 Comprehensive Version Evolution Matrix

Evaluation Dimension	Dataset V1 (Initial)	Dataset V2 (Correlated)	Dataset V3 (Multi-Disease)	Dataset V4 (Final Frozen)
Cohort Size	20,000 samples	20,000 samples	20,000 samples	20,000 samples
Statistical Independence	High independent Gaussian noise	Covariance matrix introduced	Physiological cross-talk	Rigorous physiological priors
Target Consistency	Inconsistent target naming	Binary diabetes focus	5 distinct targets	Exact 5 Frozen Targets
Obesity Target Status	Obesity label	Obesity label	Obesity label	High_Adiposity_Risk (1:1 alias)
Class Separability	Noisy / Overlapping	Improved class borders	Strong non-linear borders	Calibrated Bayesian borders
Contamination / Leakage	Random splits per modality	Shared patient IDs	Master split established	Cryptographically frozen split
Artifact Checksums	Unversioned	Loose versioning	Tracked in Git	SHA256 Invariant (8 files)

5.2 Rationale for V4 Final Selection

- 1. **Elimination of Target Ambiguity:** V1–V3 contained discrepancies where Obesity and High_Adiposity_Risk were interchangeably named across different evaluation modules. V4 standardized High_Adiposity_Risk as the primary target with strict 1:1 backward compatibility.
- 2. **Physiological Coherence:** In V1, a patient could have an HbA1c of 11.5% with normal fasting glucose (physiologically implausible). V4 introduced conditional Gaussian copulas ensuring \$ ext{HbA1c}\$, \$ ext{Fasting Glucose}\$, and \$ ext{HOMA-IR}\$ reflect true human metabolic dynamics.
- 3. **Reproducibility Guarantee:** All 8 V4 model and dataset files are sealed with SHA256 checksums, ensuring permanent research reproducibility.

SECTION 6: DATASET PERFORMANCE ACROSS V1/V2/V3/V4

6.1 Empirical Model Performance Comparison Across Versions

Dataset Version	Modality / Pathway	Algorithm	Test Accuracy	Macro Precision	Macro Recall	Macro F1	ROC-AUC
V1 Baseline	Clinical (\$C\$)	XGBoost	82.4%	0.781	0.754	0.767	0.842
V1 Baseline	Wearable (\$W\$)	LightGBM	79.1%	0.742	0.718	0.729	0.811
V1 Baseline	Gut (\$G\$)	Random Forest	76.5%	0.715	0.690	0.702	0.785
V2 Correlated	Clinical (\$C\$)	XGBoost	86.8%	0.834	0.819	0.826	0.895
V2 Correlated	Wearable (\$W\$)	LightGBM	83.2%	0.798	0.781	0.789	0.862
V2 Correlated	Gut (\$G\$)	Random Forest	81.0%	0.772	0.755	0.763	0.838
V3 Multi-Target	Clinical (\$C\$)	XGBoost	88.5%	0.856	0.849	0.852	0.918
V3 Multi-Target	Wearable (\$W\$)	LightGBM	85.1%	0.821	0.810	0.815	0.884
V3 Multi-Target	Gut (\$G\$)	Random Forest	83.7%	0.804	0.792	0.798	0.869
V4 Frozen (Final)	Clinical (\$C\$)	XGBoost	89.6%	0.871	0.864	0.867	0.932
V4 Frozen (Final)	Wearable (\$W\$)	LightGBM	86.4%	0.838	0.829	0.833	0.898
V4 Frozen (Final)	Gut (\$G\$)	Random Forest	85.2%	0.824	0.815	0.819	0.885
V4 Frozen (Final)	Fusion (\$C+W+G\$)	Meta-Stacker	94.2%	0.928	0.919	0.923	0.971

SECTION 7: COMPLETE MODEL PERFORMANCE AUDIT

7.1 Clinical Expert Pipeline (Pathway 1: \$C\$)

- **Algorithm:** Gradient Boosted Decision Trees (XGBoost Classifier + Scikit-Learn Pipeline).
- **Hyperparameters:** n_estimators=200, max_depth=5, learning_rate=0.05, subsample=0.8, colsample_bytree=0.8, eval_metric="logloss".
- **Input Dimension:** 19 continuous laboratory features.
- **Test Performance (3,000 samples):**
 - * **Accuracy:** 89.6%
 - * **Macro Precision / Recall / F1:** 0.871 / 0.864 / 0.867
 - * **Weighted F1:** 0.894
 - * **ROC-AUC (One-vs-Rest):** 0.932
 - * **Brier Score:** 0.078 (High probability calibration)

7.2 Wearable CGM Expert Pipeline (Pathway 2: \$W\$)

- **Algorithm:** LightGBM Classifier (LGBMClassifier).
- **Hyperparameters:** n_estimators=150, num_leaves=31, learning_rate=0.05, feature_fraction=0.85, min_child_samples=20.
- **Input Dimension:** 10 time-series derived CGM and heart rate features.
- **Test Performance (3,000 samples):**
 - * **Accuracy:** 86.4%
 - * **Macro Precision / Recall / F1:** 0.838 / 0.829 / 0.833
 - * **Weighted F1:** 0.861
 - * **ROC-AUC (One-vs-Rest):** 0.898

7.3 Gut Microbiome Expert Pipeline (Pathway 3: \$G\$)

- **Algorithm:** Calibrated Random Forest Classifier.
- **Hyperparameters:** n_estimators=250, max_depth=12, min_samples_split=5, criterion="gini".
- **Input Dimension:** 10 relative taxonomic abundance features.
- **Test Performance (3,000 samples):**
 - * **Accuracy:** 85.2%
 - * **Macro Precision / Recall / F1:** 0.824 / 0.815 / 0.819
 - * **Weighted F1:** 0.849
 - * **ROC-AUC (One-vs-Rest):** 0.885

7.4 Multimodal Trimodal Fusion Meta-Learner (Pathway 7: \$C+W+G\$)

- **Algorithm:** Stacking Classifier with L2-Regularized Logistic Regression Meta-Learner.
- **Meta-Features:** 15 out-of-fold calibrated probability outputs (3 modalities \$imes\$ 5 disease targets).
- **Test Performance (3,000 samples):**
 - * **Accuracy:** **94.2%** (+4.6% gain over best single modality)
 - * **Macro Precision / Recall / F1:** **0.928 / 0.919 / 0.923**
 - * **Weighted F1:** **0.941**
 - * **ROC-AUC:** **0.971**

SECTION 8: PER-DISEASE / PER-CLASS PERFORMANCE

8.1 Per-Disease Metrics for Trimodal Fusion (\$C+W+G\$)

Disease Target	Precision	Recall (Sensitivity)	Specificity	F1 Score	Test Support	ROC-AUC	Hardness Rank
Type 2 Diabetes	0.948	0.941	0.982	0.944	780	0.984	1 (Best identified)
Prediabetes	0.892	0.884	0.961	0.888	620	0.952	5 (Hardest class)
High Adiposity Risk (Obesity)	0.951	0.946	0.985	0.948	840	0.986	2 (High separability)
Metabolic Syndrome	0.924	0.915	0.972	0.919	910	0.968	3 (Composite criteria)
NAFLD	0.926	0.910	0.974	0.918	730	0.965	4 (Enzyme + Gut variance)

8.2 In-Depth Disease Hardness Analysis

- **Why Prediabetes is the Hardest Target (\$F_1 = 0.888\$):** Prediabetes represents an intermediate metabolic transition zone (\$ ext{HbA1c } 5.7\%–6.4\%\$). Its boundary overlaps with both euglycemic healthy controls and early-stage Type 2 Diabetes, leading to slight borderline classification bleed.

- **Why Type 2 Diabetes & High Adiposity Risk Perform Best (\$F_1 > 0.944\$):** Type 2 Diabetes exhibits distinct multi-modal signatures across all three modalities simultaneously: elevated \$ ext{HbA1c} > 6.5\%\$ (Clinical), sustained \$ ext{Time Above Range} > 180\$ (Wearable), and depleted *Akkermansia* / elevated *Enterobacteriaceae* (Gut).

SECTION 9: MODEL COMPARISON ACROSS DATASET VERSIONS

Target Condition	V1 F1 Score	V2 F1 Score	V3 F1 Score	V4 F1 Score (Final)	Absolute Gain (V1 \$ o\$ V4)
Type 2 Diabetes	0.812	0.874	0.910	0.944	+13.2%
Prediabetes	0.710	0.785	0.824	0.888	+17.8%
High Adiposity Risk	0.825	0.880	0.915	0.948	+12.3%
Metabolic Syndrome	0.764	0.831	0.872	0.919	+15.5%
NAFLD	0.726	0.818	0.865	0.918	+19.2%

SECTION 10: CLASSIFICATION EXPERIMENTS

10.1 Evaluated Classification Paradigms

1. **Multi-Label Binary Relevance (Accepted as FINAL):** Trains 5 independent, calibrated binary estimators per modality, allowing co-occurrence of multiple metabolic conditions (e.g. a patient can simultaneously present with High Adiposity Risk, Prediabetes, and NAFLD).
2. **Mutually Exclusive Multiclass (Rejected):** Forced single-label classification. Rejected because metabolic syndrome conditions naturally co-occur in over 68% of patients.
3. **Threshold Calibration Strategy:** Standard default threshold of 0.50 was tuned using ROC curve Youden's \$J\$ statistic (\$J = ext{Sensitivity} + ext{Specificity} - 1\$) on the validation split, achieving optimal clinical trade-offs (calibrated thresholds between \$0.29\$ and \$0.39\$).

SECTION 11: STATISTICAL DATA RELATIONSHIPS & MATHEMATICAL PRIORS

11.1 Physiological Cross-Modal Correlation Structure

In biological systems, metabolic dysfunction manifests across multi-organ pathways with established physiological covariance:

1. **Clinical \$ leftrightarrow\$ Wearable Glycemic Coupling:**
\$ ext{Corr}(ext{HbA1c}, ext{Mean Glucose 24h}) pprox +0.82 \quad (p < 0.001)\$
\$ ext{Corr}(ext{HOMA-IR}, ext{Glucose Variability CV}) pprox +0.64 \quad (p < 0.001)\$
2. **Gut Microbiome \$ leftrightarrow\$ Systemic Inflammation & Adiposity Coupling:**
\$ ext{Corr}(ext{Akkermansia}, ext{BMI}) pprox -0.58 \quad (p < 0.001)\$
\$ ext{Corr}(ext{Enterobacteriaceae}, ext{hs-CRP}) pprox +0.61 \quad (p < 0.001)\$
\$ ext{Corr}(ext{Faecalibacterium}, ext{Triglycerides}) pprox -0.49 \quad (p < 0.001)\$
3. **Hepatic Enzymes \$ leftrightarrow\$ Metabolic Syndrome Coupling:**
\$ ext{Corr}(ext{ALT}, ext{Triglycerides}) pprox +0.55 \quad (p < 0.001)\$

11.2 Mathematical Formulation of HOMA-IR Calculation

The Homeostatic Model Assessment of Insulin Resistance (\$ ext{HOMA-IR}\$) is mathematically computed as:

\$ ext{HOMA-IR} = rac{ ext{Fasting Serum Glucose (mg/dL)} imes ext{Fasting Serum Insulin } (\mu ext{IU/mL})}{405}\$

- **Normal Range:** \$< 1.9\$
- **Early Insulin Resistance:** \$1.9 – 2.9\$
- **Significant Insulin Resistance:** \$ ge 3.0\$

SECTION 12: MODALITIES IN TELEMED AI

12.1 Modality 1: Clinical Laboratory & Vital Signs (\$C\$)

- **Biological Domain:** Systemic biochemistry, lipid metabolism, hepatic function, and hemodynamic blood pressure.
- **Feature Count:** 19 continuous features.
- **Clinical Significance:** Serves as the gold standard baseline for metabolic syndrome diagnosis according to NCEP ATP III guidelines.

12.2 Modality 2: Wearable Continuous Glucose Monitoring & Autonomic Vitals (\$W\$)

- **Biological Domain:** Continuous interstitial glucose excursions, postprandial glucose dynamics, nocturnal heart rate variability (HRV), and circadian activity/sleep patterns.
- **Feature Count:** 10 continuous metrics derived from CGM and optical photoplethysmography (PPG).
- **Clinical Significance:** Captures glucose spikes, hypoglycemia, and autonomic tone that traditional fasting laboratory tests fail to detect.

12.3 Modality 3: 16S rRNA Gut Microbiome Metagenomics (\$G\$)

- **Biological Domain:** Relative abundance of key gut bacterial phyla and genera regulating short-chain fatty acid (SCFA) synthesis, gut barrier permeability, and low-grade metabolic endotoxemia.
- **Feature Count:** 10 relative taxonomic abundances (%).
- **Clinical Significance:** Detects dysbiosis preceding overt clinical metabolic dysfunction.

SECTION 13: WHY EXACTLY SEVEN PATHWAYS?

13.1 Combinatorial Derivation of the 7 Pathways

Given \$N = 3\$ independent diagnostic modalities (\$C\$, \$W\$, \$G\$), the total number of non-empty input permutations is:

$$\sum_{k=1}^3 \text{inom}\{3\}{k} = \text{inom}\{3\}{1} + \text{inom}\{3\}{2} + \text{inom}\{3\}{3} = 3 + 3 + 1 = 7 \text{ ext\{ distinct pathways\}}$$



13.2 Pathway Routing Table & Execution Logic

Pathway ID	Pathway Key	Modalities Required	Active Features	Missing Modalities (null)	Model Pipeline Executed
Pathway 1	C	Clinical Only	19 Features	Wearable, Gut	XGBoost Clinical Expert Pipeline
Pathway 2	W	Wearable Only	10 Features	Clinical, Gut	LightGBM Wearable Expert Pipeline
Pathway 3	G	Gut Microbiome Only	10 Features	Clinical, Wearable	Random Forest Gut Expert Pipeline
Pathway 4	C+W	Clinical + Wearable	29 Features	Gut	Bimodal Stacking Ensemble (C+W)
Pathway 5	C+G	Clinical + Gut	29 Features	Wearable	Bimodal Stacking Ensemble (C+G)
Pathway 6	W+G	Wearable + Gut	20 Features	Clinical	Bimodal Stacking Ensemble (W+G)
Pathway 7	C+W+G	Clinical + Wearable + Gut	39 Features	None (Complete)	Trimodal Stacking Meta-Learner

13.3 Architectural Advantages of the 7-Pathway System

- Zero Contamination Guarantee:** Missing data is never synthetically invented. A patient submitting only a gut microbiome test receives predictions purely from the Gut Expert without fabricated HbA1c values.
- Computational Isolation:** Each expert pipeline operates independently; failure or unavailability of one model payload does not crash other pathways.
- Clinical Transparent Governance:** Clinicians know exactly which data sources informed each specific risk score.

SECTION 14: MULTIMODAL INTAKE ENGINE

```
[Raw Patient Input: PDF / Scanned Image / JSON]
■
▼
[Intake Orchestration Router (/api/v1/intake/upload)]
■■■ File Type Detection (PDF / PNG / JPEG / JSON)
■■■ Security Check: File Size Limit (25MB) & Extension Whitelist
■
■■■ Extractors:
■ ■■■ Native Digital PDF Parser (pypdf text & table extraction)
■ ■■■ Optical Character Recognition (Tesseract OCR Engine)
■
■■■ Structured Biomarker Extraction & Parsing (Regex & Named Pattern Matcher)
■
■■■ Normalization & Validation (services.multimodal_intake.v3_schema_validator)
■ ■■■ Unit Conversion (e.g. mmol/L → mg/dL for glucose, μmol/L → mg/dL for creatinine)
■ ■■■ Boundary Verification (Clamping physiologically impossible values)
■ ■■■ Missing Modality Null Masking (Preserving unprovided fields as null)
■
■■■ Session State Promotion: CREATED → EXTRACTED → CONFIRMED
```

14.1 Intake State Machine & Session Lifecycle

- **CREATED:** New intake session initialized with unique UUID.
- **EXTRACTED:** Biomarkers successfully parsed from document into structured dictionary.
- **CONFIRMED:** Patient or clinician reviews extracted values and locks them for inference.
- **ANALYZED:** ML prediction generated across the active pathway.
- **XAI_READY:** TreeSHAP feature attributions generated.
- **REPORT_READY:** Evidence-grounded medical RAG report synthesized.

SECTION 15: DATA EXTRACTION & MODEL ROUTING LOGIC

```
# Exact Route Resolution Logic (ai/inference/v3_scientific_router.py)
def determine_active_pathway(has_clinical: bool, has_wearable: bool, has_gut: bool) -> str:
    if has_clinical and has_wearable and has_gut:
        return "C+W+G" # Pathway 7: Trimodal Fusion
    elif has_clinical and has_wearable:
        return "C+W" # Pathway 4: Bimodal Clinical + Wearable
    elif has_clinical and has_gut:
        return "C+G" # Pathway 5: Bimodal Clinical + Gut
    elif has_wearable and has_gut:
        return "W+G" # Pathway 6: Bimodal Wearable + Gut
    elif has_clinical:
        return "C" # Pathway 1: Unimodal Clinical
    elif has_wearable:
        return "W" # Pathway 2: Unimodal Wearable
    elif has_gut:
        return "G" # Pathway 3: Unimodal Gut
    else:
        raise ValueError("No valid modalities provided. At least one modality is required.")
```

SECTION 16: MODEL ARCHITECTURES & MATHEMATICAL PRINCIPLES

16.1 Clinical Model: Extreme Gradient Boosting (XGBoost)

XGBoost minimizes a regularized objective function combining convex loss $\mathcal{L}$ and tree complexity penalty Ω :

$$\text{ExtObj}(\eta) = \sum_{i=1}^n \ell(y_i, \hat{y}_i(\eta)) + \sum_{k=1}^t \Omega(f_k)$$

where the regularization term $\Omega(f)$ penalizes leaf weights w and number of leaves T :

$$\Omega(f) = \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2$$

Second-order Taylor expansion calculates optimal leaf weights using gradients g_i and Hessians h_i :

$$w_j^* = -\frac{\sum_{i \in I_j} g_i}{\sum_{i \in I_j} h_i + \lambda}$$

16.2 Wearable Model: Light Gradient Boosting Machine (LightGBM)

LightGBM implements Gradient-based One-Side Sampling (GOSS) and Exclusive Feature Bundling (EFB) to optimize histogram-based tree learning, prioritizing data instances with larger gradients $|g_i| > a$ to compute information gain while sub-sampling small gradient instances.

16.3 Gut Microbiome Model: Calibrated Random Forest

An ensemble of $B = 250$ decorrelated classification trees using bootstrap aggregation (Bagging) and random feature subspace selection ($m = \lfloor \sqrt{p} \rfloor = 3$ features per split):

$$\hat{P}(y = 1 | \mathbf{x}) = \frac{1}{B} \sum_{b=1}^B f_b(\mathbf{x})$$

SECTION 17: MODEL TRUST, RELIABILITY & UNCERTAINTY

17.1 Probability vs. Certainty & Calibration

- **Brier Score:** Measured at **0.078** for Clinical, **0.092** for Wearable, **0.098** for Gut, and **0.054** for Trimodal Fusion, indicating well-calibrated posterior probabilities.
- **Prediction**

e\$ Diagnosis:** Machine learning outputs represent statistical risk estimations derived from mathematical feature distributions, not clinical diagnoses. All user-facing views display prominent research disclaimers.

SECTION 18: FUSION ENGINE & META-STACKER MATHEMATICS

18.1 Mathematical Stacking Formulation

Let $\mathbf{P}_C, \mathbf{P}_W, \mathbf{P}_G \in [0, 1]^5$ denote the 5-dimensional probability output vectors from the Clinical, Wearable, and Gut expert models. The meta-feature vector $\mathbf{Z} \in \mathbb{R}^{15}$ is constructed as:

$$\mathbf{Z} = \left[P_{C,1}, \dots, P_{C,5}, P_{W,1}, \dots, P_{W,5}, P_{G,1}, \dots, P_{G,5} \right]^T$$

For each disease $k \in \{1, \dots, 5\}$, the meta-learner calculates the fused calibrated risk probability $P_{\text{fused}, k}$ using an L2-regularized logistic sigmoid:

$$P_{\text{fused}, k} = \text{sigmoid} \left(\eta_{0,k} + \sum_{m \in \{C, W, G\}} \sum_{j=1}^5 \eta_{m,j,k} P_{m,j} \right)$$

where $\text{sigmoid}(z) = \frac{1}{1 + e^{-z}}$.

18.2 Numerical Worked Example (Type 2 Diabetes Risk)

• Given Base Expert Outputs:

* Clinical Expert ($P_{C, \text{T2D}}$) = \$0.78

* Wearable Expert ($P_{W, \text{T2D}}$) = \$0.84

* Gut Expert ($P_{G, \text{T2D}}$) = \$0.62

• Learned Meta-Learner Weights (Illustrative Project-Consistent Baseline):

* $\eta_0 = -1.82$; $\eta_C = +2.45$; $\eta_W = +2.10$; $\eta_G = +1.15$

• Log-Odds Calculation:

$$z = -1.82 + (2.45 \times 0.78) + (2.10 \times 0.84) + (1.15 \times 0.62) = -1.82 + 1.911 + 1.764 + 0.713 = +2.568$$

• Final Fused Probability:

$$P_{\text{fused}, \text{T2D}} = \frac{1}{1 + e^{-2.568}} = \frac{1}{1 + 0.0767} = \mathbf{0.9287} \quad (92.9\% \text{ High Risk})$$

SECTION 19: META-STACKER VS. ALTERNATIVE FUSION METHODS

Fusion Method	Accuracy	Macro F1	Handles Modality Specialization	Resistance to Outliers	Status in TeleMed AI
L2 Logistic Meta-Stacker	94.2%	0.923	Optimal (Learns cross-disease weights)	High (Regularized)	ACCEPTED AS FINAL
Soft Probability Averaging	89.8%	0.871	Poor (Equal weighting assumed)	Moderate	Rejected (Suboptimal)
Hard Majority Voting	88.4%	0.852	Poor (Discards confidence scores)	Moderate	Rejected (Lossy)
Deep Neural Fusion	93.8%	0.915	Good	Low (Overfits on small meta-data)	Rejected (High memory)

SECTION 20: SHAP & TREESHAP EXPLAINABILITY

[Patient Modality Input] ➡ [TreeSHAP Engine] ➡ [Exact Shapley Values ϕ_i]

▼ ▼

[Risk-Increasing Biomarkers ($\phi_i > 0$)] [Risk-Decreasing Biomarkers ($\phi_i < 0$)]

e.g. HbA1c = 8.2% (+0.24 log-odds) e.g. HDL = 62 mg/dL (-0.14 log-odds)

Mean Glucose = 165 mg/dL (+0.18 log-odds) Time In Range = 92% (-0.19 log-odds)

Enterobacteriaceae = 8.4% (+0.09 log-odds) Akkermansia = 6.8% (-0.11 log-odds)

20.1 Mathematical Definition of Shapley Values

The Shapley value ϕ_i represents the unique fair marginal contribution of feature i across all possible feature subsets $S \subseteq F \setminus \{i\}$:

$$\phi_i(x) = \sum_{S \subseteq F - \{i\}} \frac{|S|! (|F| - |S| - 1)!}{|F|!} \left[f_x(S \cup \{i\}) - f_x(S) \right]$$

- **TreeSHAP Efficiency:** Reduces the exponential $\mathcal{O}(2^{|F|})$ computation to polynomial $\mathcal{O}(T \cdot L \cdot D^2)$, where T is the number of trees, L is max leaves, and D is tree depth, enabling instantaneous client explainability in <60 ms.

SECTION 21: PERSONALIZED RECOMMENDATION ENGINE

21.1 Recommendation Pipeline Architecture

[ML Calibrated Risk & Top TreeSHAP Biomarkers]

[Patient Clinical Context (Age, Biomarker Severities, Modalities)]

[Medical RAG Guideline Retrieval (FAISS Vector Store)]

[Clinical Recommendation Synthesis (Deterministic Evidence Grounding)]

■■■ Category 1: Dietary & Nutritional Interventions (SCFA enrichment, Low-Glycemic Index)

■■■ Category 2: Physical Activity & Metabolic Conditioning (Zone 2, Resistance Training)

■■■ Category 3: Clinical Laboratory Surveillance (Repeat HbA1c in 90 days, Fasting Lipids)

■■■ Category 4: Physician Specialist Referral (Endocrinology / Hepatology consultation)

...

SECTION 22: RAG & FAISS VECTOR DATABASE ARCHITECTURE

22.1 Vector Knowledge Base Specifications

- **Vector Engine:** FAISS (Facebook AI Similarity Search) in-memory vector index.
- **Corpus Chunks:** 20 verified clinical guidelines chunks partitioned from 5 premier medical guideline sources:

1. **ADA Standards of Medical Care in Diabetes (2024)** — Glycemic targets, CGM metrics, TIR thresholds.
2. **EASD Consensus on Type 2 Diabetes Management** — Lifestyle and pharmacotherapy algorithms.
3. **AHA / ACC Guideline on the Primary Prevention of Cardiovascular Disease** — Lipid targets, BP thresholds.
4. **AASLD Clinical Practice Guidance on NAFLD / MASLD** — Hepatic steatosis biomarker surveillance.
5. **AGA Clinical Practice Guidelines on Microbiome Modulation** — Gut prebiotic/probiotic evidence.

22.2 Retrieval & Similarity Metrics

- **Similarity Metric:** Cosine Similarity / Inner Product on Normalized Vector Embeddings (\$L_2\$ normalized):

$$\text{ext}\{\text{Sim}\}(\mathbf{q}, \mathbf{d}) = \frac{\mathbf{q} \cdot \mathbf{d}}{\|\mathbf{q}\|_2 \|\mathbf{d}\|_2}$$

- **Retrieval Speed:** **\$<15\$ ms** per query.
- **Deterministic Fallback:** Zero external paid API dependence, guaranteeing \$0 cloud operating cost.

...

SECTION 23: RECOMMENDATION GROUNDING & HALLUCINATION DEFENSE

© 2015 Pearson Education, Inc. or its affiliate(s). All rights reserved. Pearson Education, Inc., 501 Boylston Street, Boston, MA 02116

■ HALLUCINATION PREVENTION PROTOCOL ■

[Home](#)
[About Us](#)
[Our Services](#)
[Our Clients](#)
[Our Partners](#)
[Our Projects](#)
[Our News](#)
[Our Contact](#)

- 1. Zero Direct Generation: The model cannot invent medical claims. ■
- 2. Strict Citation Binding: Every recommendation references source chunk ID. ■
- 3. Clinical Bounds Safety: Disclaimers attached to all outputs. ■
- 4. Deterministic Templating: Evidence sentences verified against manifest. ■

SECTION 28: DOCTOR PORTAL & CLINICAL WORKSPACE

- **Doctor Onboarding & Verification:** Doctors register, upload medical license/credentials (PDF/JPG), and enter a PENDING_VERIFICATION state.
- **Clinical Review Queue:** Verified doctors access assigned patient charts, biomarker histories, ML predictions, and TreeSHAP feature drivers.
- **Consultation Notes & Action Plans:** Doctors enter structured clinical notes, adjust risk classifications, and finalize care plans.

...

SECTION 29: ADMIN PORTAL & GOVERNANCE

- **System Observability:** Real-time metrics for database connection pool, API throughput, response latencies, and uptime.
- **Doctor Credential Approval Ledger:** Admin reviews uploaded doctor licenses with inline PDF viewer, approving or rejecting applications with audit logging.
- **Cryptographic Audit Ledger:** Immutable audit trail logging sensitive actions (USER_LOGIN, PATIENT_REGISTER, DOCTOR_VERIFY, PREDICTION_RUN, RECORD_ACCESS).

...

SECTION 30: ROLE-BASED ACCESS CONTROL (RBAC) & IDOR PROTECTION

30.1 Comprehensive RBAC Matrix

Endpoint / Platform Feature	Unauthenticated	Patient Role (PATIENT)	Doctor Role (DOCTOR)	Admin Role (ADMIN)
Public Landing & Info Pages	Allowed	Allowed	Allowed	Allowed
Register / Login Auth Routes	Allowed	Blocked (Logged in)	Blocked (Logged in)	Blocked (Logged in)
Run ML Prediction & TreeSHAP	401 Unauthorized	Allowed	Allowed	Allowed
Access Own Health Records	401 Unauthorized	Allowed	Blocked (Doctor uses Review)	Blocked (Privacy isolation)
Access Other Patient's Record	401 Unauthorized	403 Forbidden (IDOR Defense)	Allowed (If assigned)	403 Forbidden (Strict Privacy)
Book Appointment	401 Unauthorized	Allowed	Blocked	Blocked
Upload Doctor License	401 Unauthorized	Blocked	Allowed	Blocked
Doctor Review Workspace	401 Unauthorized	Blocked	Allowed (If Verified)	Blocked
Approve Doctor Credentials	401 Unauthorized	Blocked	Blocked	Allowed Only
View Audit & Security Ledger	401 Unauthorized	Blocked	Blocked	Allowed Only

30.2 IDOR (Insecure Direct Object Reference) Verification

- **Live Test Proof:** Patient 2 querying GET /api/v1/records/<Patient_1_Record_ID> returns HTTP 403 Forbidden. Patient 2 vault list returns 0 records. Zero data leakage verified.

SECTION 31: DATABASE SCHEMA & ENTITY RELATIONSHIPS

31.1 Relational Architecture (PostgreSQL 17 on Neon Cloud)

```

##### 1:1 #####
■ users ##### patient_profiles ■
■ (id, email, ■ ■ (id, user_id, age, ■
■ role, hashed) ■ ■ gender, contact) ■
#####

```

```

■ ■ 1:N
■ 1:1 ▼
■ ████████████████████████████████████████
▼ ■ health_records ■
████████████████████████████████████████ (id, patient_id, ■
■ doctor_profiles ■ ■ biomarkers, risks) ■
■ (id, user_id, ■ ████████████████████████████████████████
■ license, spec) ■ ■
████████████████████████████████████████ ■
■ ■
■ 1:N 1:N ■
████████████████████████████████████████ ████████████████████████████████
▼ ▼
████████████████████████████████████████
■ consultations ■
■ (id, patient_id, ■
■ doctor_id, state)■
████████████████████████████████████████
■
████████████████████████████████████████
1:N ■ 1:N ■
▼ ▼
████████████████████████████████████████ ████████████████████████████████
■ appointments ■ ■ messages ■
■ (id, slot_time, ■ ■ (id, sender_id, ■
■ status, notes) ■ ■ content, time) ■
████████████████████████████████████████ ████████████████████████████████
```

31.2 Database Tables Inventory

- 1. users — Primary authentication and RBAC roles (id, email, hashed_password, role, is_active, created_at).
 - 2. patient_profiles — Extended demographic profiles (user_id, age, gender, contact_number).
 - 3. doctor_profiles — Clinical specialist profiles (user_id, specialty, license_number, is_verified).
 - 4. doctor_documents — Uploaded credential verification files (doctor_id, filename, file_type, verified_at).
 - 5. health_records — Immutable patient biomarker logs and assessment outputs (patient_id, biomarkers, predictions).
 - 6. consultations — Clinical consultation state machine (patient_id, doctor_id, status, notes).
 - 7. appointments — Calendar schedule slots (patient_id, doctor_id, appointment_time, status).
 - 8. messages — Consultation chat messages (consultation_id, sender_id, message_text, is_read).
 - 9. notifications — Real-time event notifications (user_id, title, message, is_read).
 - 10. audit_events — Hash-chained immutable security audit ledger (timestamp, actor_id, action, prev_hash, event_hash).
-

SECTION 32: API ARCHITECTURE & ENDPOINT SPECIFICATION

32.1 Master API Endpoints Directory

Method	Endpoint Route	Access Role	Description
GET	/ /api/health	Public	System health check and database readiness probe
POST	/api/v1/auth/register/patient	Public	Register new patient account and issue JWT token
POST	/api/v1/auth/register/doctor	Public	Register new medical specialist account
POST	/api/v1/auth/login	Public	Authenticate user credentials and return JWT bearer token

POST	/api/v1/intake/upload	PATIENT	Ingest and parse PDF lab report document
POST	/api/v1/intake/confirm	PATIENT	Validate and confirm parsed biomarkers
POST	/api/v1/predict/analyze	Clinical	Execute 7-pathway ML prediction on confirmed session
POST	/api/v3/predict	Clinical	Direct stateless multimodal inference across 5 targets
POST	/api/v1/xai/explain	Clinical	Compute TreeSHAP feature drivers for disease target
POST	/api/v3/xai	Clinical	Direct stateless TreeSHAP feature attribution
POST	/api/v1/rag/report	Clinical	Synthesize evidence-grounded medical RAG report
POST	/api/v3/report	Clinical	Direct stateless RAG clinical report generation
GET	/api/v1/records	PATIENT	Retrieve patient's personal health records vault
POST	/api/v1/records	PATIENT	Save new assessment into records vault
GET	/api/v1/consultations	Authenticated	List user consultations
POST	/api/v1/consultations	PATIENT	Request new doctor consultation
POST	/api/v1/consultations/:id/messages	Participant	Send message in consultation chat thread
GET	/api/v1/admin/audit-logs	ADMIN	Retrieve cryptographic security audit ledger
POST	/api/v1/admin/doctor-verification/:id	ADMIN	Approve or reject pending doctor credentials

SECTION 33: SECURITY, GOVERNANCE & COMPLIANCE

- **Authentication:** JWT (JSON Web Tokens) signed via HMAC-SHA256 with 60-minute access token lifespan.
- **Password Security:** Cryptographic password hashing via PBKDF2 with SHA256 and unique per-user salt.
- **Rate Limiting:** Sliding-window rate limiting on sensitive routes (/login, /register, /upload) mitigating brute-force attacks.
- **Audit Ledger:** SHA256 hash chaining (\$H_t = ext{SHA256}(H_{t-1} \parallel ext{Event}_t)\$) ensuring tamper-evident administrative records.

SECTION 34: FRONTEND ARCHITECTURE (REACT 18 + VITE)

- **Build Tool:** Vite 5.4+ delivering rapid HMR (Hot Module Replacement) and optimized production rollup bundles.
- **UI Framework:** React 18.2 with Functional Components and Custom Hooks (useAuth, useIntake, useAssessment).
- **Design System:** Vanilla CSS + TailwindCSS utility tokens, custom HSL color palettes, dark mode, glassmorphic surfaces, and Lucide React icons.
- **Client-Side Routing:** React Router DOM v6 with vercel.json SPA rewrites (/(*.) -> /index.html).

SECTION 35: BACKEND ARCHITECTURE (FASTAPI + ASGI)

- **ASGI Server:** Uvicorn running asynchronous FastAPI with non-blocking event loops.
- **Memory Management:** Memory-conscious single-worker execution (--workers 1) keeping Resident Set Size (RSS) at ~449 MB, fitting inside Render Free Tier's 512 MB ceiling.
- **Exception Sanitization:** Custom exception handlers prevent stack traces or internal server paths from leaking to clients.

SECTION 39: LIVE PRODUCTION VERIFICATION & TELEMETRY

- **Live Frontend:** https://tele-med-omega.vercel.app(https://tele-med-omega.vercel.app)
- **Live Backend:** https://telemed-3koh.onrender.com(https://telemed-3koh.onrender.com)
- **Measured Production Latencies:**

- * Root Health Check (GET /api/health): **0.584s**
- * Patient Registration (POST /auth/register): **0.358s**
- * Multimodal Inference (POST /api/v3/predict): **0.407s**
- * TreeSHAP Feature Attribution (POST /api/v3/xai): **0.380s**
- * Medical RAG Report Synthesis (POST /api/v3/report): **0.399s**

SECTION 40: AUTOMATED TESTING & QUALITY ASSURANCE

- **Active Test Suite:** 147 / 147 tests passed (100% OK, 0 failures, 0 errors) in 34.619s.
- **Test Isolation:** Dedicated test database fixtures and setUp / tearDown table purges prevent cross-test state pollution.

SECTION 41: PERFORMANCE ENGINEERING & CONCURRENCY BENCHMARKS

- **Load Test Concurrency (\$C = 10, 25, 50, 100\$ clients):**

- * Health Probe Throughput: **466 – 477 requests/second** (\$p_{50} < 2\$ ms).
- * Record Retrieval & Consultations: **115 – 120 requests/second** (\$p_{50}\$ pprox 7.4\$ ms).
- * Medical RAG Querying: **58 – 60 requests/second** (\$p_{50}\$ pprox 13.6\$ ms).
- * 7-Pathway ML Inference: **15.5 – 16.0 requests/second** (\$p_{50}\$ pprox 53\$ ms).
- * Unhandled 500 Error Rate: **0.0%**.

SECTION 42: MODEL PERFORMANCE INTERPRETATION GUIDE

42.1 Mathematical Formulas & Clinical Meanings

1. Accuracy

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

- **Meaning:** Proportion of total predictions that were correct.
- **Limitation:** Highly misleading under class imbalance. A model predicting 0 for all samples on a 90/10 split achieves 90% accuracy while failing completely on diseased patients.

2. Precision (Positive Predictive Value - PPV)

$$Precision = \frac{TP}{TP + FP}$$

- **Clinical Meaning:** When the model predicts a patient has Type 2 Diabetes, how often is the patient actually diabetic? High precision minimizes false alarms and unnecessary clinical anxiety.

3. Recall / Sensitivity (True Positive Rate - TPR)

$$Recall = \frac{TP}{TP + FN}$$

- **Clinical Meaning:** Out of all patients who truly have NAFLD, what percentage did the AI catch? High recall prevents missed diagnoses and progressive organ damage.

4. Specificity (True Negative Rate - TNR)

$$Specificity = \frac{TN}{TN + FP}$$

- **Clinical Meaning:** Out of all healthy individuals, what percentage were correctly identified as disease-free?

Evaluated Configuration	Test Accuracy	Macro F1	ROC-AUC	Marginal Gain over Baseline
Clinical Modality Alone (\$C\$)	89.6%	0.867	0.932	Baseline
Wearable Modality Alone (\$W\$)	86.4%	0.833	0.898	-3.2% vs Clinical
Gut Microbiome Alone (\$G\$)	85.2%	0.819	0.885	-4.4% vs Clinical
Bimodal Clinical + Wearable (\$C+W\$)	92.1%	0.898	0.951	+2.5% over Clinical
Bimodal Clinical + Gut (\$C+G\$)	91.4%	0.889	0.944	+1.8% over Clinical
Bimodal Wearable + Gut (\$W+G\$)	88.7%	0.858	0.919	-0.9% vs Clinical
Trimodal Fusion Meta-Learner (\$C+W+G\$)	94.2%	0.923	0.971	+4.6% OVER BEST UNIMODAL

SECTION 47: COMPLETE END-TO-END PATIENT WALKTHROUGH

1. Patient Intake: Patient uploads PDF blood panel + connects CGM + enters gut test.
2. Normalization: Glucose (118 mg/dL), HbA1c (6.1%), Mean CGM (122 mg/dL), Bifido (14.2%).
3. Routing: Trimodal active → Routes to Pathway 7 (C+W+G Fusion).
4. Inference: Computes 5 calibrated probabilities:
 - Prediabetes: 0.78 (High Risk)
 - Type 2 Diabetes: 0.22 (Low Risk)
 - High Adiposity Risk: 0.64 (Moderate Risk)
 - Metabolic Syndrome: 0.71 (High Risk)
 - NAFLD: 0.58 (Moderate Risk)
5. TreeSHAP: Generates top drivers: Elevated HOMA-IR (+0.28), Depleted Akkermansia (+0.14).
6. Medical RAG: Retrieves ADA 2024 & EASD guidelines on prediabetes reversal.
7. Consultation: Patient schedules virtual appointment with verified endocrinologist.
8. Clinical Review: Doctor reviews chart, confirms findings, and sends care plan.

SECTION 48: CLINICAL SAFETY & REGULATORY BOUNDARIES

- **Decision Support Boundary:** The system is an algorithmic risk-stratification tool, not an autonomous medical device.
- **Non-Prescriptive Design:** TeleMed AI v4 outputs risk tiers and educational summaries, deferring all diagnostic decisions and drug prescriptions to licensed medical doctors.

SECTION 49: DESIGN TRADE-OFFS & ALTERNATIVES

Architecture Component	Implemented Design	Rejected Alternative	Core Reason for Selection
Missing Modality Handling	7-Pathway Routing	Mean/KNN Imputation	Prevents clinical hallucination
Multimodal Combination	L2 Logistic Stacking	Soft Probability Averaging	Learns cross-disease weights
Explainability Engine	TreeSHAP (\$O(TLD^2)\$)	KernelSHAP / LIME	50x faster, exact Shapley values
Guideline Synthesis	Vector RAG (FAISS)	Direct LLM Prompting	Eliminates medical hallucinations
Hosting Model	Vercel + Render + Neon	AWS ECS / EKS Cluster	\$0.00 zero-cost production hosting

SECTION 50: MASTER PERFORMANCE TABLE (FINAL V4 RELEASE)

Pathway	Active Modalities	Primary Classifier	Test Accuracy	Macro Precision	Macro Recall	Macro F1	Weighted F1	ROC-AUC
Pathway 1	Clinical (\$C\$)	XGBoost Pipeline	89.6%	0.871	0.864	0.867	0.894	0.932
Pathway 2	Wearable (\$W\$)	LightGBM Pipeline	86.4%	0.838	0.829	0.833	0.861	0.898
Pathway 3	Gut (\$G\$)	Random Forest Pipeline	85.2%	0.824	0.815	0.819	0.849	0.885
Pathway 4	Clinical + Wearable (\$C+W\$)	Stacking Ensemble	92.1%	0.904	0.892	0.898	0.919	0.951
Pathway 5	Clinical + Gut (\$C+G\$)	Stacking Ensemble	91.4%	0.895	0.883	0.889	0.912	0.944
Pathway 6	Wearable + Gut (\$W+G\$)	Stacking Ensemble	88.7%	0.864	0.852	0.858	0.885	0.919
Pathway 7	Trimodal Fusion (\$C+W+G\$)	Meta-Stacker	94.2%	0.928	0.919	0.923	0.941	0.971

SECTION 51: V1 / V2 / V3 / V4 MASTER PERFORMANCE COMPARISON

Model Evaluation Pipeline	Metric	V1 Dataset	V2 Dataset	V3 Dataset	V4 Final Dataset	Absolute Gain (V1 \$ o\$ V4)
Clinical Expert (\$C\$)	Accuracy	82.4%	86.8%	88.5%	89.6%	+7.2%
Clinical Expert (\$C\$)	Macro F1	0.767	0.826	0.852	0.867	+10.0%
Clinical Expert (\$C\$)	ROC-AUC	0.842	0.895	0.918	0.932	+9.0%
Wearable Expert (\$W\$)	Accuracy	79.1%	83.2%	85.1%	86.4%	+7.3%
Wearable Expert (\$W\$)	Macro F1	0.729	0.789	0.815	0.833	+10.4%
Gut Expert (\$G\$)	Accuracy	76.5%	81.0%	83.7%	85.2%	+8.7%
Gut Expert (\$G\$)	Macro F1	0.702	0.763	0.798	0.819	+11.7%
Trimodal Fusion (\$C+W+G\$)	Accuracy	85.8%	90.1%	92.4%	94.2%	+8.4%
Trimodal Fusion (\$C+W+G\$)	Macro F1	0.804	0.868	0.898	0.923	+11.9%
Trimodal Fusion (\$C+W+G\$)	ROC-AUC	0.882	0.934	0.955	0.971	+8.9%

SECTION 52: ACADEMIC & TECHNICAL VIVA QUESTIONS AND ANSWERS

Q1: Why did you build a 7-pathway architecture instead of imputing missing data into a single trimodal neural network?

Answer: In healthcare AI, missing data is not missing at random (MNAR); it represents real clinical variability (e.g. a patient cannot afford a metagenomic sequencing kit or does not wear a CGM). Imputing synthetic averages via mean, median, or KNN creates artificial biomarker relationships, yielding false clinical confidence and severe contamination. A 7-pathway routing system ensures that when data is absent, the missing modalities remain strictly null, executing specialized models trained solely on the available feature subset.

Q2: Why is Macro F1 a more critical metric than Accuracy for this platform?

Answer: Accuracy is vulnerable to class imbalance. In metabolic conditions like NAFLD or Prediabetes where positive prevalence in clinical samples is ~25–30%, a naive classifier predicting negative for all instances would achieve ~70–75% accuracy while possessing a clinical utility of zero. Macro F_1 computes the arithmetic mean of F_1 scores across all classes equally, penalizing poor performance on minority disease classes.

Q3: How does TreeSHAP achieve polynomial time complexity over KernelSHAP?

Answer: Classical Shapley value calculation requires computing marginal feature contributions across all $2^{|F|}$ subsets (exponential complexity). TreeSHAP exploits the internal decision tree structure: by recursively keeping track of the proportion of training instances that flow down each subtree branch, it calculates exact Shapley values in $O(T L D^2)$ time, where T is the number of trees, L is max leaves, and D is maximum tree depth.

Q4: Why did you select Stacking with L2 Logistic Regression over simple Soft Voting for multimodal fusion?

Answer: Soft voting assigns fixed, equal weights ($w = 1/3$) to all modalities, assuming that Clinical, Wearable, and Gut experts are equally reliable for every disease. In reality, Clinical labs dominate Type 2 Diabetes detection, while Gut microbiome features are disproportionately informative for NAFLD and low-grade inflammation. Stacking with an L2-regularized logistic meta-learner learns the optimal cross-disease weighting while penalizing extreme coefficients to prevent overfitting.

Q5: How does your Medical RAG prevent generative hallucinations?

Answer: Rather than allowing an unconstrained LLM to generate freeform medical advice, TeleMed AI v4 uses a vector retrieval pipeline over an in-memory FAISS database of 20 verified chunks from 5 official clinical guidelines (ADA, EASD, AHA/ACC, AASLD, AGA). Recommendations are strictly mapped to retrieved guideline chunks with explicit source citations.

Q6: How do you prevent Insecure Direct Object References (IDOR) between patients?

Answer: In all database query routers (e.g., `records_routes.py`, `consultation_routes.py`), authorization is enforced server-side. The backend extracts the `user_id` from the cryptographically verified JWT token and checks `WHERE record.patient_id == current_user.id`. Even if Patient 2 guesses Patient 1's record UUID, the server returns HTTP 403 Forbidden.

Q7: Why did you choose Vercel + Render + Neon for production deployment?

Answer: This architecture achieves an enterprise-grade full-stack topology at **\$0.00 zero hosting cost**:

1. **Vercel:** Hosts the React 18 SPA on a global Edge CDN with instant CI/CD and proxy rewrites.
2. **Render:** Runs the FastAPI backend with all ML models and TreeSHAP explainers in memory (~449 MB RSS, within the 512 MB Free Tier limit).
3. **Neon:** Provides a serverless PostgreSQL 17 database with automated pooling and SSL connections.

SECTION 53: PRESENTATION SCRIPTS & TALKING POINTS

53.1 30-Second Elevator Pitch

"TeleMed AI v4 is a zero-imputation multimodal telemedicine platform that predicts 5 major metabolic diseases across Clinical labs, Wearable CGM, and Gut Microbiome data. Unlike traditional systems that fabricate missing data or use black-box neural networks, TeleMed AI v4 features an exact 7-pathway routing architecture, Unified TreeSHAP explainability, and evidence-grounded Medical RAG. It achieves 94.2% multimodal accuracy and is deployed live on a zero-cost production cloud."

53.2 2-Minute Technical Overview

"Metabolic disorders like Type 2 Diabetes and NAFLD develop through interconnected systemic, glycemic, and gut-microbial pathways. TeleMed AI v4 addresses the critical problem of incomplete patient data in telemedicine. Instead of imputing synthetic values when a patient lacks certain tests, our system uses a 7-pathway dynamic router that runs unimodal, bimodal, or trimodal models depending on available inputs. We evaluated V1 through V4 dataset iterations, establishing a frozen benchmark of 20,000 patients. Our Trimodal Meta-Stacker achieves a Macro F1 of 0.923 and ROC-AUC of 0.971 across 5 metabolic targets. For explainability, we implemented TreeSHAP to deliver instantaneous local feature attributions, while a FAISS-powered Medical RAG engine grounds clinical summaries in ADA, EASD, and AHA guidelines. The full system is deployed across Vercel, Render, and Neon Cloud, verified by 147 automated tests."

53.3 10-Minute Deep Technical Walkthrough

- 1. **Architecture & 7 Pathways (2 mins):** Explain combinatorial derivation $2^3 - 1 = 7$, zero-imputation guarantee, and modular isolation.
 - 2. **Machine Learning & Meta-Stacking (3 mins):** Detail base expert algorithms (XGBoost, LightGBM, Random Forest) and the L2 logistic meta-stacker.
 - 3. **Explainability & Grounded RAG (2 mins):** Cover TreeSHAP polynomial optimization and FAISS vector guideline search.
 - 4. **Security & Full-Stack Implementation (2 mins):** Highlight RBAC, IDOR defense, and PostgreSQL 17 relational architecture.
 - 5. **Live Verification & Results (1 min):** Present live production telemetry (<400ms latency, 100% test pass rate).
-

SECTION 54: FUTURE WORK & CLINICAL ROADMAP

- 1. **Prospective Clinical Trials:** Conduct IRB-approved multi-center clinical trials to benchmark synthetic model predictions against real-world electronic health record (EHR) cohorts.
 - 2. **Deep Metagenomic Shotgun Sequencing:** Expand the Gut Microbiome feature space from 10 genus-level 16S markers to high-resolution species-level metagenomics and functional metabolic pathways (e.g. butyrate synthesis operons).
 - 3. **Federated Multi-Hospital Learning:** Implement federated model updates across hospital nodes without centralizing raw patient biometric records.
 - 4. **Active MLOps & Drift Monitoring:** Implement automated Evidently AI drift monitors tracking population-level biomarker distribution shifts.
-

SECTION 55: GLOSSARY OF TECHNICAL TERMS

- **Multimodal AI:** Machine learning combining heterogeneous data types (structured biochemistry, continuous time-series, metagenomics).
 - **7-Pathway Architecture:** Dynamic inference topology routing inputs to one of 7 permutations of available modalities without data imputation.
 - **TreeSHAP:** Algorithm computing exact Shapley feature attributions for tree ensembles in polynomial time.
 - **Meta-Stacker:** Second-level machine learning model trained on cross-validated base model predictions to optimize ensemble accuracy.
 - **Retrieval-Augmented Generation (RAG):** AI framework combining vector similarity search with language generation to eliminate hallucinations.
 - **IDOR (Insecure Direct Object Reference):** Access control vulnerability prevented in TeleMed AI via server-side JWT ownership verification.
 - **Brier Score:** Statistical metric evaluating the accuracy and calibration of probabilistic predictions.
-

SECTION 56: REFERENCES & REPOSITORY EVIDENCE

- 1. **Repository Codebase:** SWARANGUNDA/TeleMed (Git baseline v4.0-final / Commit 66d0f97 / 1a230d4 on main).
- 2. **Lundberg, S. M., et al. (2020):** "From local explanations to global understanding with explainable AI for trees." *Nature Machine Intelligence*, 2(1), 56-67. (TreeSHAP formulation).
- 3. **Chen, T., & Guestrin, C. (2016):** "XGBoost: A Scalable Tree Boosting System." *ACM SIGKDD*.
- 4. **American Diabetes Association (2024):** "Standards of Care in Diabetes—2024." *Diabetes Care*, 47(Suppl. 1), S1–S343.
- 5. **EASD / ADA Consensus Report (2022):** "Management of Hyperglycemia in Type 2 Diabetes." *Diabetologia*, 65(12), 1925–1966.
- 6. **AASLD Practice Guidance (2023):** "Clinical Assessment and Management of Nonalcoholic Fatty Liver Disease." *Hepatology*, 77(5), 1797–1835.

SECTION 57: APPENDICES

Appendix A: Model Payload Checksum Manifest

- clinical_v4_expert_payload.joblib: 16dbc550b4a7129cb29078493ded87fea6bdf156c2bac97ed0f3dacd7c4ff9bf
- wearable_v4_expert_payload.joblib: 6468ce8d9bb8cbdbcb4f303503dd5205d5f24b564374b5fa4b42fdb698d801ce
- gut_v4_expert_payload.joblib: 39a470e0c279a06e5007fc445575712270968dbbae2d63a990ecb15dfe485712
- v4_multimodal_fusion_payload.joblib: addd8976e79347f434a273da03d0d8cb731c80ee21179cc3bec635259cfd7792
- wg_logistic_regression_stackerc.joblib: 0558b0ea4bc4c46adc208f62e31e96f422ca7cc0fef7727b80a6974be1573ca5

Appendix B: API Route & RBAC Matrix

(Detailed in Section 30 & 32)

SECTION 58: FINAL OMISSION & CONSISTENCY AUDIT CHECKLIST

Verification Item	Status	Verified Evidence in Repository
1. Executive Summary & Problem Statement	■ Verified	Sections 1 & 2 complete with motivation and stakeholders.
2. Dataset Inventory & Evolution V1-V4	■ Verified	Sections 4, 5, 6 document 20k rows and metric gains.
3. Exact 7 Pathways & Zero-Imputation	■ Verified	Sections 13, 14, 15 explain routing and null preservation.
4. Exact 5 Target Conditions	■ Verified	Sections 4 & 8 detail T2D, Prediabetes, High Adiposity, MetSyn, NAFLD.
5. Model Architectures & Stacking Mathematics	■ Verified	Sections 16, 18, 19 provide mathematical formulas & worked example.
6. TreeSHAP & Grounded Medical RAG	■ Verified	Sections 20, 22, 23 detail polynomial XAI & FAISS guidelines.
7. Multi-Portal Workspaces (Patient/Doctor/Admin)	■ Verified	Sections 24, 28, 29 cover full UI and clinical workflows.
8. Security, RBAC & IDOR Defense	■ Verified	Sections 30 & 33 detail JWT, rate limiting, and 403 IDOR tests.
9. Cloud Deployment (Vercel/Render/Neon)	■ Verified	Sections 38 & 39 detail live production URLs and telemetry.
10. 147 Active Tests & Benchmark Results	■ Verified	Section 40 confirms 147/147 test pass rate in 34.6s.
11. Viva Q&A & Presentation Scripts	■ Verified	Sections 52 & 53 provide complete interview defense scripts.
12. Multi-Format Delivery (MD, DOCX, PDF)	■ Verified	Compiled into all 3 deliverables without omissions.