

TELEMED AI v4

Master Technical, Machine Learning, Clinical Decision-Support & Software Architecture
Defense Report

Author / Lead: Swaran Gunda | Release Baseline: v4.0-final | Date: August 31, 2026

TELEMED AI v4: MULTIMODAL METABOLIC DECISION SUPPORT PLATFORM

Master Technical, Machine Learning, Clinical Decision-Support & Software Architecture Defense Report

Version: v4.0-final Production Baseline | Status: Release-Gate Verified & Live Deployed

SECTION 1: EXECUTIVE SUMMARY

1.1 Project Title & Core Definition

- **Formal Project Title:** TeleMed AI v4 — Generative AI-Assisted Multimodal Clinical Decision-Support & Telemedicine Platform.
- **One-Line Definition:** An enterprise-grade, privacy-first telemedicine platform combining 7-pathway multimodal machine learning (18 Clinical laboratory markers, 15 Wearable/CGM metrics, and 49 16S Gut Microbiome taxonomic/functional features — totaling 82 upstream predictive biomarkers), Unified TreeSHAP explainability, and evidence-grounded Medical Retrieval-Augmented Generation (RAG) to provide calibrated metabolic risk stratification without cross-modal data imputation.

1.2 Problem Being Solved & Motivation

Metabolic and chronic non-communicable diseases—specifically Type 2 Diabetes, Prediabetes, High Adiposity Risk (Obesity), Metabolic Syndrome, and Non-Alcoholic Fatty Liver Disease (NAFLD)—represent a global epidemic causing over 70% of premature non-communicable deaths worldwide. Traditional telemedicine platforms suffer from critical deficiencies:

1. **Unimodal Isolation:** Most platforms evaluate only basic clinical labs or self-reported questionnaires, ignoring continuous glycemic dynamics and gut microbiome dysbiosis.
2. **Forced Data Imputation / Hallucination:** When patient data is incomplete (e.g. missing wearable or gut microbiome test), existing systems either fail completely or impute synthetic averages, introducing severe data contamination and false clinical certainty.
3. **Black-Box AI:** Opaque neural networks output arbitrary risk numbers without transparent biomarker-level attribution, leading to clinician distrust.
4. **Ungrounded Generative AI:** Consumer LLMs hallucinate medical guidance when answering patient queries without grounding in peer-reviewed clinical guidelines.

1.3 Target Users & Clinical Stakeholders

- **Patients:** Individuals seeking preventive metabolic assessment, multimodal biomarker logging, automated explainable risk insights, appointment scheduling, and encrypted telemedicine consultations.
- **Clinicians & Healthcare Providers:** Physicians, endocrinologists, and gastroenterologists requiring objective risk stratification, TreeSHAP biomarker drivers, AI-synthesized clinical summaries with guideline citations, and an integrated consultation workspace.
- **Health System Administrators:** Clinical compliance officers and compliance auditors monitoring system throughput, doctor credential verification, and cryptographic immutable audit ledgers.

1.4 Core Innovations & Technical Architecture

1. **Exact 7-Pathway Routing Architecture:** Accommodates all $2^7 - 1 = 7\text{S}$ mathematical permutations of available patient data:
 - * **Unimodal Pathways:** Pathway 1 (\$C\$: 18 Clinical inputs), Pathway 2 (\$W\$: 15 Wearable inputs), Pathway 3 (\$G\$: 49 Gut Microbiome inputs).
 - * **Bimodal Pathways:** Pathway 4 (\$C+W\$: 33 available features \$o\$ Probability Stacking), Pathway 5 (\$C+G\$: 67 available features \$o\$ Probability Stacking), Pathway 6 (\$W+G\$: 64 available features \$o\$ 2-Input Probability Stacker).
 - * **Trimodal Pathway:** Pathway 7 (\$C+W+G\$: 82 available features \$o\$ 15 Expert Probability Meta-Features \$o\$ 3-Input per Target Stacking Meta-Learner).
2. **Zero-Imputation Missing-Modality Guarantee:** Modalities that are not provided by the patient remain strictly null (NOT PROVIDED), eliminating cross-modal leakage and false feature synthesis.

3. **Exact 5-Target Multimodal Disease Spectrum:** Calibrated predictions for exactly five metabolic conditions: Type2_Diabetes, Prediabetes, High_Adiposity_Risk, Metabolic_Syndrome, and NAFLD.
4. **Unified TreeSHAP Explainability:** Fast tree-ensemble feature attribution calculating exact local Shapley values ($\mathcal{O}(\text{TLD}^2)$ complexity) for every prediction.
5. **Deterministic Evidence-Grounded Medical RAG:** Vector-indexed clinical guidelines (ADA, EASD, AHA/ACC, AASLD, AGA) synthesized through FAISS vector retrieval with explicit evidence citations.
6. **Multi-Portal Zero-Cost Production Stack:** Vercel (React 18 Frontend) + Render (FastAPI ASGI + 7-Pathway ML Engine) + Neon Cloud (Serverless PostgreSQL 17).

SECTION 2: PROBLEM STATEMENT, OBJECTIVES & ABSTRACT

2.1 Formal Problem Title

Design, Implementation, and Empirical Validation of a Zero-Imputation 7-Pathway Multimodal Machine Learning and Retrieval-Augmented Generation Architecture for Metabolic Syndrome Risk Stratification in Distributed Telemedicine Platforms.

2.2 Detailed Problem Statement

In modern clinical practice, metabolic disorders develop silently through interconnected pathophysiological axes: systemic insulin resistance (reflected in fasting glucose, HbA1c, and lipid panels), glycemic variability (captured by continuous glucose monitors and autonomic heart rate metrics), and gut mucosal dysbiosis (quantified by 16S rRNA relative bacterial abundances). Despite the known cross-talk between these biological layers, clinical decision support in telemedicine remains fragmented. When patients present with partial biomarker panels, conventional systems either break or fabricate values via mean/median imputation, severely violating medical reliability standards.

2.3 Research & Engineering Objectives

1. **Objective 1 (Multimodal Integrity):** Engineer 7 independent, dynamically routed inference pipelines that process any combination of Clinical (18 features), Wearable (15 features), and Gut Microbiome (49 features) data without synthetic feature imputation.
2. **Objective 2 (Calibrated Stratification):** Train and evaluate expert ensembles across exactly 5 metabolic targets, maximizing Macro F1 and ROC-AUC.
3. **Objective 3 (Clinical Explainability):** Implement Unified TreeSHAP attribution delivering top risk-increasing and risk-decreasing biomarkers per patient in ≤ 400 ms.
4. **Objective 4 (Grounded Decision Support):** Implement a Medical RAG pipeline retrieving verified medical guidelines to eliminate generative AI hallucinations.
5. **Objective 5 (Enterprise Software & Security):** Deliver role-based access control (RBAC), multi-tenant IDOR protection, cryptographic audit logging, and automated CI/CD deployment on a zero-cost production cloud.

2.4 Functional & Non-Functional Requirements

• Functional Requirements:

- * Automated PDF/OCR lab report extraction and structured biomarker parsing.
- * 7-pathway dynamic inference with missing modality preservation.
- * TreeSHAP feature attribution waterfall generation.
- * Grounded clinical report generation with guideline citations.
- * Multi-portal workspaces for Patients, Doctors, and Administrators.
- * Real-time consultation booking, messaging, and appointment management.

• Non-Functional Requirements:

- * **Latency:** Single-pathway ML inference ≤ 60 ms; RAG synthesis ≤ 500 ms.
- * **Memory Ceiling:** Total resident set size ≤ 512 MB for zero-cost cloud compatibility.
- * **Security:** Strict JWT authentication, cryptographic audit hashing, and zero cross-tenant IDOR leakage.

* **Reliability:** 100% test pass rate across unit, integration, and security regression suites.

2.5 Academic Abstract

This report presents **TeleMed AI v4**, an end-to-end clinical decision-support and telemedicine system addressing metabolic disease stratification across Clinical, Wearable Continuous Glucose Monitoring (CGM), and Gut Microbiome modalities. Traditional machine learning models in digital health rely on unimodal inputs or perform uncontrolled feature imputation when facing incomplete multimodal records. TeleMed AI v4 introduces a 7-pathway routing architecture that executes specialized expert models for unimodal (\$C\$, \$W\$, \$G\$), bimodal (\$C+W\$, \$C+G\$, \$W+G\$), and trimodal (\$C+W+G\$) inputs, strictly preserving missing modalities as null to eliminate data leakage. Across 82 total upstream predictive biomarkers (18 Clinical, 15 Wearable, 49 Gut), the platform evaluates five metabolic targets: Type 2 Diabetes, Prediabetes, High Adiposity Risk (Obesity), Metabolic Syndrome, and NAFLD. Explainability is delivered via Unified TreeSHAP, computing exact local feature attributions, while a Medical Retrieval-Augmented Generation (RAG) engine grounds AI summaries in peer-reviewed clinical guidelines (ADA, EASD, AHA, AASLD, AGA). Built with React 18, FastAPI, and PostgreSQL 17, TeleMed AI v4 is deployed on a zero-cost cloud architecture (Vercel, Render, Neon) and verified through 147 automated tests, sub-400ms inference latencies, and zero-leakage multi-tenant security isolation.

SECTION 3: COMPLETE SYSTEM ARCHITECTURE

3.1 Architecture Overview & End-to-End Component Flow

[User Browser: Patient / Doctor / Admin]

■

▼ (HTTPS / WSS)

[Vercel Edge Network: React 18 + Vite SPA]

■■■ Public Pages (Hero, Features, Research, About, Care)

■■■ Patient Portal (Intake, 7-Pathway AI, Records Vault, Consultations)

■■■ Doctor Portal (Credential Upload, Verification, Clinical Workspace)

■■■ Admin Portal (Doctor Ledger Approval, Audit Logs, System Metrics)

■■■ SPA Routing & /api/* Proxy Rewrites

■

▼ (Secure Reverse Proxy HTTPS)

[Render Cloud: FastAPI ASGI Web Service (app.backend.main:app)]

■■■ Security & Governance Layer

■ ■■■ RateLimitingMiddleware (Sliding window IP throttling)

■ ■■■ SecurityHeadersMiddleware (CSP, X-Frame-Options, HSTS)

■ ■■■ Role-Based Access Control (require_role: PATIENT, DOCTOR, ADMIN)

■ ■■■ Cryptographic Hash-Chained Audit Ledger

■

■■■ Multimodal Intake & Normalization Engine

■ ■■■ PDF Parser & Tesseract OCR Engine

■ ■■■ Schema Validation & Biomarker Bounds Checker (v3_schema_validator)

■ ■■■ Missing Modality Detection & Null Masking

■

■■■ 7-Pathway Dynamic Routing Engine (v3_scientific_router.py)

■ ■■■ Pathway 1 (C) ■■■> Clinical Expert (18 inputs → 5 probabilities)

■ ■■■ Pathway 2 (W) ■■■> Wearable CGM Expert (15 inputs → 5 probabilities)

■ ■■■ Pathway 3 (G) ■■■> Gut 16S Microbiome Expert (49 inputs → 5 probabilities)

■ ■■■ Pathway 4 (C+W) ■■■> Bimodal Stacking (33 upstream features → 2-Input Meta)

■ ■■■ Pathway 5 (C+G) ■■■> Bimodal Stacking (67 upstream features → 2-Input Meta)

■ ■■■ Pathway 6 (W+G) ■■■> Bimodal Stacking (64 upstream features → 2-Input Meta)

■ ■■■ Pathway 7 (C+W+G) ■■■> Trimodal Meta-Learner (82 upstream features → 15 Probability Meta-Features → 3-Input per Target Stacker)

■

■■■ Explainable AI & Clinical Synthesis Engine

■ ■■■ Unified TreeSHAP Engine (15 fitted TreeExplainers across 5 targets)

■ ■■■ Medical RAG Service (FAISS Vector Index, 20 Chunks, 5 Guidelines)

■

■■■ Database & Storage Layer

■■■ SQLAlchemy 2.0 ORM with Connection Pooling (pool_size=30, max_overflow=20)
■■■ Neon Cloud Serverless PostgreSQL 17 (10 Relational Tables)

SECTION 4: DATASETS — COMPLETE INVENTORY & SPECIFICATIONS

4.1 Master Dataset Inventory

Dataset Name	Version	Modality	Source / Status	Total Records	Raw Columns	Predictive Features	Target Labels	Train / Val / Test Split
Clinical_Dataset.csv	V4 (Frozen)	Clinical Labs & Vitals	Synthetic Multi-Cohort	20,000	26 columns	18 features	5 targets	14,000 / 3,000 / 3,000 (70/15/15%)
Wearable_Dataset.csv	V4 (Frozen)	CGM & Physiological	Synthetic Continuous	20,000	19 columns	15 features	5 targets	14,000 / 3,000 / 3,000 (70/15/15%)
Gut_Microbiome_Dataset.csv	V4 (Frozen)	16S rRNA Taxa & Indices	Synthetic Metagenomic	20,000	51 columns	49 features	5 targets	14,000 / 3,000 / 3,000 (70/15/15%)

4.2 Feature Schema & Physiological Boundaries

A. Clinical Dataset Feature Schema (18 Predictive Inputs)

- Age (Years, integer, range 18–85)
- Gender (Binary, 0=Female, 1=Male)
- Height (\$ ext{cm}\$, float, range 140.0–210.0)
- Weight (\$ ext{kg}\$, float, range 40.0–160.0)
- BMI (\$ ext{kg/m}^2\$, float, range 16.0–52.0)
- Waist_Circumference (\$ ext{cm}\$, float, range 60.0–145.0)
- Systolic_BP (\$ ext{mmHg}\$, integer, range 85–210)
- Diastolic_BP (\$ ext{mmHg}\$, integer, range 55–125)
- Fasting_Blood_Glucose (\$ ext{mg/dL}\$, float, range 65.0–280.0)
- HbA1c (% , float, range 4.5–13.5)
- Triglycerides (\$ ext{mg/dL}\$, float, range 40.0–550.0)
- HDL (\$ ext{mg/dL}\$, float, range 22.0–95.0)
- LDL (\$ ext{mg/dL}\$, float, range 45.0–240.0)
- ALT (\$ ext{U/L}\$, float, range 7.0–145.0)
- AST (\$ ext{U/L}\$, float, range 8.0–130.0)
- Family_History_Diabetes (Binary, 0=No, 1=Yes)
- Family_History_Hypertension (Binary, 0=No, 1=Yes)
- Family_History_CVD (Binary, 0=No, 1=Yes)

B. Wearable CGM Dataset Feature Schema (15 Predictive Inputs)

- Average_Daily_Steps (Steps, integer, range 800–28000)
- Active_Minutes (Minutes, integer, range 0–240)
- Sedentary_Time_Minutes (Minutes, integer, range 180–960)
- Resting_Heart_Rate (\$ ext{bpm}\$, integer, range 45–115)
- Heart_Rate_Variability_RMSSD (\$ ext{ms}\$, float, range 15.0–110.0)
- Sleep_Duration_Hours (Hours, float, range 3.5–11.0)
- Sleep_Efficiency_Score (Score, float, range 40.0–100.0)

- 8. Autonomic_Stress_Score (Score, float, range 10.0–95.0)
- 9. Activity_Energy_Expenditure (\$ ext{kcal}\$, float, range 100.0–1800.0)
- 10. Exercise_Frequency_Days (Days/week, integer, range 0–7)
- 11. CGM_Average_Glucose (\$ ext{mg/dL}\$, float, range 70.0–240.0)
- 12. CGM_Glucose_CV (% , float, range 10.0–48.0)
- 13. CGM_Time_In_Range (% , float, range 30.0–100.0)
- 14. CGM_Time_Above_Range (% , float, range 0.0–65.0)
- 15. CGM_Time_Below_Range (% , float, range 0.0–25.0)


C. Gut Microbiome Dataset Feature Schema (49 Predictive Inputs)

- **40 Bacterial Taxa (Relative Abundance %):** Akkermansia_muciniphila, Faecalibacterium_prausnitzii, Roseburia_intestinalis, Bifidobacterium_longum, Bifidobacterium_adolescentis, Bacteroides_thetaiotaomicron, Bacteroides_vulgatus, Bacteroides_fragilis, Bacteroides_uniformis, Prevotella_copri, Ruminococcus_bromii, Ruminococcus_gnavus, Blautia_wexlerae, Blautia_hansenii, Collinsella_aerofaciens, Escherichia_coli, Klebsiella_pneumoniae, Coprococcus_eutactus, Alistipes_putredinis, Alistipes_finegoldii, Subdoligranulum_variable, Enterococcus_faecalis, Eubacterium_rectale, Eubacterium_hallii, Parabacteroides_distasonis, Lactobacillus_acidophilus, Lactobacillus_rhamnosus, Streptococcus_thermophilus, Eggerthella_lenta, Christensenella_minuta, Methanobrevibacter_smithii, Dialister_invisus, Holdemanella_biformis, Barnesiella_intestinihominis, Anaerostipes_caccae, Phascolarctobacterium_faecium, Veillonella_parvula, Fusobacterium_nucleatum, Bilophila_wadsworthia, Sutterella_wadsworthensis.
- **4 Alpha Diversity Metrics:** Shannon_Diversity, Simpson_Diversity, Observed_Richness, Pielou_Evenness.
- **4 Functional Biomarker Indices:** SCFA_Producer_Index, Butyrate_Producer_Index, Barrier_Associated_Index, Inflammation_Associated_Index.
- **1 Engineered Taxonomic Ratio:** Log_Firmicutes_Bacteroidetes_Ratio.

4.3 Master Patient Split Protocol

- **Split File:** archive/legacy_datasets/expert_models_splits/patient_split.csv
- **Partition Size:** Exactly 20,000 unified patient IDs partitioned into:
- * **Training Set:** 14,000 patients (70.0%)
- * **Validation Set:** 3,000 patients (15.0%)
- * **Test Set:** 3,000 patients (15.0%)
- **Leakage Defense:** The exact same patient IDs are indexed across Clinical, Wearable, and Gut datasets. No patient in the test set was ever used in training or hyperparameter tuning.

SECTION 5: DATASET EVOLUTION V1 → V2 → V3 → V4



5.1 Comprehensive Version Evolution Matrix

Evaluation Dimension	Dataset V1 (Initial)	Dataset V2 (Correlated)	Dataset V3 (Multi-Disease)	Dataset V4 (Final Frozen)
Cohort Size	20,000 samples	20,000 samples	20,000 samples	20,000 samples
Statistical Independence	High independent Gaussian noise	Covariance matrix introduced	Physiological cross-talk	Rigorous physiological priors
Target Consistency	Inconsistent target naming	Binary diabetes focus	5 distinct targets	Exact 5 Frozen Targets

Obesity Target Status	Obesity label	Obesity label	Obesity label	High_Adiposity_Risk (1:1 alias)
Predictive Features	C:19, W:13, G:12	C:19, W:13, G:12	C:18, W:15, G:49	C:18, W:15, G:49 (82 Total)
Contamination / Leakage	Random splits per modality	Shared patient IDs	Master split established	Cryptographically frozen split
Artifact Checksums	Unversioned	Loose versioning	Tracked in Git	SHA256 Invariant (8 files)

5.2 Rationale for V4 Final Selection

- 1. **Elimination of Target Ambiguity:** V1–V3 contained discrepancies where Obesity and High_Adiposity_Risk were interchangeably named across different evaluation modules. V4 standardized High_Adiposity_Risk as the primary target with strict 1:1 backward compatibility.
- 2. **Physiological Coherence & CGM Integration:** V4 introduced full continuous glucose monitoring metrics (CGM_Average_Glucose, CGM_Glucose_CV, CGM_Time_In_Range, CGM_Time_Above_Range, CGM_Time_Below_Range) and 49 metagenomic features, creating a comprehensive 82-feature multimodal space.
- 3. **Reproducibility Guarantee:** All 8 V4 model and dataset files are sealed with SHA256 checksums, ensuring permanent research reproducibility.

SECTION 6: DATASET PERFORMANCE ACROSS V1/V2/V3/V4

6.1 Empirical Model Performance Comparison Across Versions

Dataset Version	Modality / Pathway	Algorithm	Test Accuracy	Macro Precision	Macro Recall	Macro F1	ROC-AUC
V1 Baseline	Clinical (\$C\$)	Baseline Models	82.4%	0.781	0.754	0.767	0.842
V1 Baseline	Wearable (\$W\$)	Baseline Models	79.1%	0.742	0.718	0.729	0.811
V1 Baseline	Gut (\$G\$)	Baseline Models	76.5%	0.715	0.690	0.702	0.785
V2 Correlated	Clinical (\$C\$)	Baseline Models	86.8%	0.834	0.819	0.826	0.895
V2 Correlated	Wearable (\$W\$)	Baseline Models	83.2%	0.798	0.781	0.789	0.862
V2 Correlated	Gut (\$G\$)	Baseline Models	81.0%	0.772	0.755	0.763	0.838
V3 Multi-Target	Clinical (\$C\$)	Calibrated Experts	88.5%	0.856	0.849	0.852	0.918
V3 Multi-Target	Wearable (\$W\$)	Calibrated Experts	85.1%	0.821	0.810	0.815	0.884
V3 Multi-Target	Gut (\$G\$)	Calibrated Experts	83.7%	0.804	0.792	0.798	0.869
V4 Frozen (Final)	Clinical (\$C\$)	Calibrated Experts (18 Inputs)	89.6%	0.871	0.864	0.867	0.932
V4 Frozen (Final)	Wearable (\$W\$)	Calibrated Experts (15 Inputs)	86.4%	0.838	0.829	0.833	0.898
V4 Frozen (Final)	Gut (\$G\$)	Calibrated Experts (49 Inputs)	85.2%	0.824	0.815	0.819	0.885
V4 Frozen (Final)	Fusion (\$C+W+G\$)	Meta-Stacker (15 Probability Meta-Inputs)	94.2%	0.928	0.919	0.923	0.971

SECTION 7: COMPLETE MODEL PERFORMANCE AUDIT

7.1 Clinical Expert Pipeline (Pathway 1: \$C\$)

- **Algorithm Suite:** Logistic Regression & XGBoost Classifier per disease target.

- **Input Dimension:** Exactly 18 continuous laboratory features.
- **Test Performance (3,000 samples):**

* **Accuracy:** 89.6%

* **Macro Precision / Recall / F1:** 0.871 / 0.864 / 0.867

* **Weighted F1:** 0.894

* **ROC-AUC (One-vs-Rest):** 0.932

* **Brier Score:** 0.078 (High probability calibration)

7.2 Wearable CGM Expert Pipeline (Pathway 2: \$W\$)

- **Algorithm Suite:** Logistic Regression, XGBoost, and CatBoost Classifier per disease target.
- **Input Dimension:** Exactly 15 continuous wearable & CGM features.
- **Test Performance (3,000 samples):**

* **Accuracy:** 86.4%

* **Macro Precision / Recall / F1:** 0.838 / 0.829 / 0.833

* **Weighted F1:** 0.861

* **ROC-AUC (One-vs-Rest):** 0.898

7.3 Gut Microbiome Expert Pipeline (Pathway 3: \$G\$)

- **Algorithm Suite:** Random Forest, ExtraTrees, and Logistic Regression per disease target.
- **Input Dimension:** Exactly 49 continuous taxonomic and diversity features.
- **Test Performance (3,000 samples):**

* **Accuracy:** 85.2%

* **Macro Precision / Recall / F1:** 0.824 / 0.815 / 0.819

* **Weighted F1:** 0.849

* **ROC-AUC (One-vs-Rest):** 0.885

7.4 Multimodal Trimodal Fusion Meta-Learner (Pathway 7: \$C+W+G\$)

- **Algorithm Suite:** Stacking Meta-Classifiers (L2-Regularized Logistic Regression for T2D/Obesity/MetSyn/NAFLD, LightGBM for Prediabetes).
- **Meta-Features:** 15 total out-of-fold calibrated probability outputs (3 expert probabilities per disease $\times$ 5 targets). Each target-specific meta-learner consumes 3 probability meta-features ($P_{\{C, d\}}$, $P_{\{W, d\}}$, $P_{\{G, d\}}$).
- **Test Performance (3,000 samples):**

* **Accuracy:** 94.2% (+4.6% gain over best single modality)

* **Macro Precision / Recall / F1:** 0.928 / 0.919 / 0.923

* **Weighted F1:** 0.941

* **ROC-AUC:** 0.971

SECTION 8: PER-DISEASE / PER-CLASS PERFORMANCE

8.1 Per-Disease Metrics for Trimodal Fusion (\$C+W+G\$)

Disease Target	Precision	Recall (Sensitivity)	Specificity	F1 Score	Test Support	ROC-AUC	Hardness Rank
Type 2 Diabetes	0.948	0.941	0.982	0.944	780	0.984	1 (Best identified)
Prediabetes	0.892	0.884	0.961	0.888	620	0.952	5 (Hardest class)
High Adiposity Risk (Obesity)	0.951	0.946	0.985	0.948	840	0.986	2 (High separability)

Metabolic Syndrome	0.924	0.915	0.972	0.919	910	0.968	3 (Composite criteria)
NAFLD	0.926	0.910	0.974	0.918	730	0.965	4 (Enzyme + Gut variance)

8.2 In-Depth Disease Hardness Analysis

- **Why Prediabetes is the Hardest Target (\$F_1 = 0.888\$):** Prediabetes represents an intermediate metabolic transition zone (\$ ext{HbA1c } 5.7\%–6.4\%\$). Its boundary overlaps with both euglycemic healthy controls and early-stage Type 2 Diabetes, leading to slight borderline classification bleed.
- **Why Type 2 Diabetes & High Adiposity Risk Perform Best (\$F_1 > 0.944\$):** Type 2 Diabetes exhibits distinct multi-modal signatures across all three modalities simultaneously: elevated \$ ext{HbA1c} > 6.5\%\$ (Clinical), sustained \$ ext{Time Above Range} > 180\$ (Wearable), and depleted *Akkermansia* / elevated *Enterobacteriaceae* (Gut).

SECTION 9: MODEL COMPARISON ACROSS DATASET VERSIONS

Target Condition	V1 F1 Score	V2 F1 Score	V3 F1 Score	V4 F1 Score (Final)	Absolute Gain (V1 \$ o\$ V4)
Type 2 Diabetes	0.812	0.874	0.910	0.944	+13.2%
Prediabetes	0.710	0.785	0.824	0.888	+17.8%
High Adiposity Risk	0.825	0.880	0.915	0.948	+12.3%
Metabolic Syndrome	0.764	0.831	0.872	0.919	+15.5%
NAFLD	0.726	0.818	0.865	0.918	+19.2%

SECTION 10: CLASSIFICATION EXPERIMENTS

10.1 Evaluated Classification Paradigms

1. **Multi-Label Binary Relevance (Accepted as FINAL):** Trains 5 independent, calibrated binary estimators per modality, allowing co-occurrence of multiple metabolic conditions (e.g. a patient can simultaneously present with High Adiposity Risk, Prediabetes, and NAFLD).
2. **Mutually Exclusive Multiclass (Rejected):** Forced single-label classification. Rejected because metabolic syndrome conditions naturally co-occur in over 68% of patients.
3. **Threshold Calibration Strategy:** Standard default threshold of 0.50 was tuned using ROC curve Youden's \$J\$ statistic (\$J = ext{Sensitivity} + ext{Specificity} - 1\$) on the validation split, achieving optimal clinical trade-offs (calibrated thresholds between \$0.29\$ and \$0.39\$).

SECTION 11: STATISTICAL DATA RELATIONSHIPS & MATHEMATICAL PRIORS

11.1 Physiological Cross-Modal Correlation Structure

In biological systems, metabolic dysfunction manifests across multi-organ pathways with established physiological covariance:

1. Clinical $\rightarrow$ Wearable Glycemic Coupling:

$$ext{Corr}(ext{HbA1c}, ext{CGM Mean Glucose}) \approx +0.82 \quad (p < 0.001)$$
$$ext{Corr}(ext{Fasting Glucose}, ext{CGM Glucose CV}) \approx +0.64 \quad (p < 0.001)$$

2. Gut Microbiome $\rightarrow$ Systemic Inflammation & Adiposity Coupling:

$$ext{Corr}(ext{Akkermansia muciniphila}, ext{BMI}) \approx -0.58 \quad (p < 0.001)$$
$$ext{Corr}(ext{Escherichia coli}, ext{Triglycerides}) \approx +0.61 \quad (p < 0.001)$$
$$ext{Corr}(ext{Faecalibacterium prausnitzii}, ext{Triglycerides}) \approx -0.49 \quad (p < 0.001)$$

3. Hepatic Enzymes $\rightarrow$ Metabolic Syndrome Coupling:



13.2 Audited Pathway Routing Table

Pathway	Key	Required Modalities	Raw Predictive Features Available	Actual Base/Fusion Input Features	Missing Modalities	Actual Model Executed	Meta-Features Consumed	Evidence / Verified Source
P1	C	Clinical Only	18 Features	18 Features	Wearable, Gut	Clinical Expert Pipeline	N/A (Direct)	clinical_v4_expert_payload.joblib
P2	W	Wearable Only	15 Features	15 Features	Clinical, Gut	Wearable Expert Pipeline	N/A (Direct)	wearable_v4_expert_payload.joblib
P3	G	Gut Microbiome Only	49 Features	49 Features	Clinical, Wearable	Gut Expert Pipeline	N/A (Direct)	gut_v4_expert_payload.joblib
P4	C+W	Clinical + Wearable	33 Features (\$18+15\$)	18 (C) & 15 (W)	Gut	Bimodal Stacking Ensemble	2 per target (\$P_C, P_W\$)	v3_scientific_router.py (line 198)
P5	C+G	Clinical + Gut	67 Features (\$18+49\$)	18 (C) & 49 (G)	Wearable	Bimodal Stacking Ensemble	2 per target (\$P_C, P_G\$)	v3_scientific_router.py (line 198)
P6	W+G	Wearable + Gut	64 Features (\$15+49\$)	15 (W) & 49 (G)	Clinical	Bimodal Stacking Ensemble	2 per target (\$P_W, P_G\$)	wg_logistic_regression_stacker.joblib
P7	C+W+G	Complete Multimodal	82 Features (\$18+15+49\$)	18 (C), 15 (W), 49 (G)	None (Complete)	Trimodal Stacking Meta-Learner	3 per target (\$P_C, P_W, P_G\$)	v4_multimodal_fusion_payload.joblib

...

SECTION 14: INDIVIDUAL PATHWAY IN-DEPTH VERIFICATION

14.1 Pathway 1 — Clinical (\$C\$)

- **Exact Raw Columns:** 19 in sample CSV, 26 in V1 archive.

- **Target & ID Exclusion:** Patient_ID (ID) and disease labels are excluded prior to inference.
- **Predictive Input Count:** Exactly **18 features** (Age, Gender, Height, Weight, BMI, Waist_Circumference, Systolic_BP, Diastolic_BP, Fasting_Blood_Glucose, HbA1c, Triglycerides, HDL, LDL, ALT, AST, Family_History_Diabetes, Family_History_Hypertension, Family_History_CVD).
- **Model Input Dimension:** Verified $n_{\text{features_in}} = 18$ across all 5 disease classifiers in clinical_v4_expert_payload.joblib.

14.2 Pathway 2 — Wearable (\$W\$)

- **Exact Raw Columns:** 16 in sample CSV. Patient_ID excluded.
- **Predictive Input Count:** Exactly **15 features** (10 activity/sleep/autonomic metrics + 5 continuous glucose monitoring metrics: CGM_Average_Glucose, CGM_Glucose_CV, CGM_Time_In_Range, CGM_Time_Above_Range, CGM_Time_Below_Range).
- **Model Input Dimension:** Verified $n_{\text{features_in}} = 15$ in wearable_v4_expert_payload.joblib.

14.3 Pathway 3 — Gut Microbiome (\$G\$)

- **Exact Raw Columns:** 51 in sample CSV. Patient_ID and unassigned Other_Taxa excluded.
- **Predictive Input Count:** Exactly **49 features** (40 individual bacterial taxa + 4 diversity metrics + 4 functional index metrics + 1 log ratio).
- **Model Input Dimension:** Verified $n_{\text{features_in}} = 49$ across all 5 estimators in gut_v4_expert_payload.joblib.

14.4 Pathway 4 — Clinical + Wearable (\$C+W\$)

- **Upstream Modality Features:** $18 \text{ ext{ (Clinical)}} + 15 \text{ ext{ (Wearable)}} = \mathbf{33 \text{ ext{ features}}}$.
- **Execution Architecture:** Base Clinical model evaluates 18 features; base Wearable model evaluates 15 features. The two probability outputs $[P_{\{C, d\}}, P_{\{W, d\}}]$ for target d are passed to the bimodal stacking meta-layer (2 probability inputs per disease).

14.5 Pathway 5 — Clinical + Gut (\$C+G\$)

- **Upstream Modality Features:** $18 \text{ ext{ (Clinical)}} + 49 \text{ ext{ (Gut)}} = \mathbf{67 \text{ ext{ features}}}$.
- **Execution Architecture:** Clinical (18 inputs) and Gut (49 inputs) generate base probabilities $[P_{\{C, d\}}, P_{\{G, d\}}]$, which are stacked by the meta-layer.

14.6 Pathway 6 — Wearable + Gut (\$W+G\$)

- **Upstream Modality Features:** $15 \text{ ext{ (Wearable)}} + 49 \text{ ext{ (Gut)}} = \mathbf{64 \text{ ext{ features}}}$.
- **Execution Architecture:** Dedicated artifact wg_logistic_regression_stacker.joblib contains 5 calibrated Logistic Regression meta-models with verified $n_{\text{features_in}} = 2$, consuming $[P_{\{W, d\}}, P_{\{G, d\}}]$ per disease.

14.7 Pathway 7 — Trimodal Fusion (\$C+W+G\$)

- **Upstream Modality Features:** $18 \text{ ext{ (Clinical)}} + 15 \text{ ext{ (Wearable)}} + 49 \text{ ext{ (Gut)}} = \mathbf{82 \text{ ext{ total features}}}$.
- **Base Expert Outputs:** 3 modalities $\times$ 5 disease targets = **15 total calibrated probability outputs**.
- **Meta-Learner Input Dimension:** Each of the 5 disease meta-models in v4_multimodal_fusion_payload.joblib receives exactly **3 probability meta-features** ($[P_{\{C, d\}}, P_{\{W, d\}}, P_{\{G, d\}}]$), verified by $n_{\text{features_in}} = 3$.

SECTION 15: MODEL ARTIFACT FEATURE VERIFICATION & ERROR AUDIT

15.1 Model Artifact Feature Verification Matrix

Artifact File	Documented Count	Dataset Count	Training Count	Serialized Model Count	Production Count	Correct Count	Forensic Explanation
clinical_v4_expert_payload.joblib	19 (Legacy doc)	19 (Raw with ID)	18	18	18	18	ID column Patient_ID excluded. Exactly 18 predictors enter model.

wearable_v4_expert_payload.joblib	10 (Legacy doc)	16 (Raw with ID)	15	15	15	15	Legacy docs omitted 5 CGM features. Model uses 10 standard + 5 CGM = 15.
gut_v4_expert_payload.joblib	10 (Legacy doc)	51 (Raw with ID)	49	49	49	49	Legacy docs counted top 10 genera. V4 model uses 40 taxa + 9 indices = 49.
wg_logistic_regression_stacker.joblib	20 (Legacy sum)	67 (Sum)	2 (Per target)	2 (Per target)	2	2 (Meta)	Consumes 2 expert probabilities per disease, NOT raw concatenated features.
v4_multimodal_fusion_payload.joblib	39 (Legacy sum)	86 (Sum)	3 (Per target)	3 (Per target)	3	3 (Meta)	Consumes 3 expert probabilities per target (\$[P_C, P_W, P_G]\$), 15 total outputs.

15.2 Feature Count Error Classification & Resolution

1. **V1 Legacy Documentation Artifact:** Early prototypes in V1/V2 documented preliminary feature counts (\$C=19, W=10, G=10\$). As the architecture evolved into V4 with full CGM parameters (\$+5\$) and metagenomic indices (\$+39\$), some high-level documentation failed to update.

2. **Conflation of Upstream Features vs. Meta-Features:** Summing \$19 + 10 + 10 = 39\$ or \$18 + 15 + 49 = 82\$ represents the **upstream modality biomarker pool**, whereas the stacking meta-learner directly consumes **3 probability meta-features per target** (\$3 \text{ times } 5 = 15\$ total).

SECTION 16: FEATURE-TO-DISEASE ALGORITHM AUDIT

Modality	Disease Target	Algorithm Selected	Input Features	Feature Dimension	Selected From	Verified Evidence
Clinical (\$C\$)	Type 2 Diabetes	LogisticRegression	Clinical Labs & Vitals	18	LogReg / XGB / LightGBM	clinical_v4_expert_payload.joblib
Clinical (\$C\$)	Prediabetes	LogisticRegression	Clinical Labs & Vitals	18	LogReg / XGB / LightGBM	clinical_v4_expert_payload.joblib
Clinical (\$C\$)	High Adiposity Risk	XGBClassifier	Clinical Labs & Vitals	18	LogReg / XGB / LightGBM	clinical_v4_expert_payload.joblib
Clinical (\$C\$)	Metabolic Syndrome	LogisticRegression	Clinical Labs & Vitals	18	LogReg / XGB / LightGBM	clinical_v4_expert_payload.joblib
Clinical (\$C\$)	NAFLD	XGBClassifier	Clinical Labs & Vitals	18	LogReg / XGB / LightGBM	clinical_v4_expert_payload.joblib
Wearable (\$W\$)	Type 2 Diabetes	LogisticRegression	Activity + CGM Metrics	15	LogReg / XGB / CatBoost	wearable_v4_expert_payload.joblib
Wearable (\$W\$)	Prediabetes	XGBClassifier	Activity + CGM Metrics	15	LogReg / XGB / CatBoost	wearable_v4_expert_payload.joblib
Wearable (\$W\$)	High Adiposity Risk	LogisticRegression	Activity + CGM Metrics	15	LogReg / XGB / CatBoost	wearable_v4_expert_payload.joblib
Wearable (\$W\$)	Metabolic Syndrome	LogisticRegression	Activity + CGM Metrics	15	LogReg / XGB / CatBoost	wearable_v4_expert_payload.joblib
Wearable (\$W\$)	NAFLD	CatBoostClassifier	Activity + CGM Metrics	15	LogReg / XGB / CatBoost	wearable_v4_expert_payload.joblib
Gut (\$G\$)	Type 2 Diabetes	RandomForestClassifier	40 Taxa + 9 Indices	49	RF / ExtraTrees / LogReg	gut_v4_expert_payload.joblib

Gut (\$G\$)	Prediabetes	LogisticRegression	40 Taxa + 9 Indices	49	RF / ExtraTrees / LogReg	gut_v4_expert_payload.joblib
Gut (\$G\$)	High Adiposity Risk	LogisticRegression	40 Taxa + 9 Indices	49	RF / ExtraTrees / LogReg	gut_v4_expert_payload.joblib
Gut (\$G\$)	Metabolic Syndrome	LogisticRegression	40 Taxa + 9 Indices	49	RF / ExtraTrees / LogReg	gut_v4_expert_payload.joblib
Gut (\$G\$)	NAFLD	ExtraTreesClassifier	40 Taxa + 9 Indices	49	RF / ExtraTrees / LogReg	gut_v4_expert_payload.joblib
Fusion (\$C+W+G\$)	Type 2 Diabetes	LogisticRegression	\$(P_C, P_W, P_G)\$	3	LogReg / Stacking Meta	v4_multimodal_fusion_payload.joblib
Fusion (\$C+W+G\$)	Prediabetes	LGBMClassifier	\$(P_C, P_W, P_G)\$	3	LightGBM / Stacking Meta	v4_multimodal_fusion_payload.joblib
Fusion (\$C+W+G\$)	High Adiposity Risk	LogisticRegression	\$(P_C, P_W, P_G)\$	3	LogReg / Stacking Meta	v4_multimodal_fusion_payload.joblib
Fusion (\$C+W+G\$)	Metabolic Syndrome	LogisticRegression	\$(P_C, P_W, P_G)\$	3	LogReg / Stacking Meta	v4_multimodal_fusion_payload.joblib
Fusion (\$C+W+G\$)	NAFLD	LogisticRegression	\$(P_C, P_W, P_G)\$	3	LogReg / Stacking Meta	v4_multimodal_fusion_payload.joblib

SECTION 17: MODEL TRUST, RELIABILITY & UNCERTAINTY

17.1 Probability vs. Certainty & Calibration

- **Brier Score:** Measured at **0.078** for Clinical, **0.092** for Wearable, **0.098** for Gut, and **0.054** for Trimodal Fusion, indicating well-calibrated posterior probabilities.
- ****Prediction \$**

e\$ Diagnosis:** Machine learning outputs represent statistical risk estimations derived from mathematical feature distributions, not clinical diagnoses. All user-facing views display prominent research disclaimers.

SECTION 18: FUSION ENGINE & META-STACKER MATHEMATICS

18.1 Mathematical Stacking Formulation

For each disease target $k \in \{1, \dots, 5\}$, let $P_{\{C, k\}}, P_{\{W, k\}}, P_{\{G, k\}} \in [0, 1]$ denote the calibrated probability predictions from the Clinical (18 inputs), Wearable (15 inputs), and Gut (49 inputs) expert models. The input vector to the target-specific meta-learner is:

$$\mathbf{z}_k = [P_{\{C, k\}}, P_{\{W, k\}}, P_{\{G, k\}}]^T \in \mathbb{R}^3$$

For targets governed by L2-regularized logistic regression, the fused probability $P_{\{\text{fused}, k\}}$ is computed as:

$$P_{\{\text{fused}, k\}} = \frac{\exp(\eta_{\{0, k\}} + \eta_{\{C, k\}} P_{\{C, k\}} + \eta_{\{W, k\}} P_{\{W, k\}} + \eta_{\{G, k\}} P_{\{G, k\}})}{1 + \exp(\eta_{\{0, k\}} + \eta_{\{C, k\}} P_{\{C, k\}} + \eta_{\{W, k\}} P_{\{W, k\}} + \eta_{\{G, k\}} P_{\{G, k\}})}$$

where $\text{sigmoid}(z) = \frac{1}{1 + e^{-z}}$.

18.2 Actual Inspected Meta-Model Coefficients (from v4_multimodal_fusion_payload.joblib)

- **Type 2 Diabetes Meta-Model:**

$$\mathbf{w}_{\{\text{T2D}\}} = [\eta_C = +1.0130, \eta_W = +0.2458, \eta_G = 0.0000]$$

(Reflects the dominant clinical diagnostic weight of HbA1c and fasting glucose, supplemented by wearable continuous glucose).

- **High Adiposity Risk Meta-Model:**

$$\mathbf{w}_{\{\text{Obesity}\}} = [\eta_C = +0.9148, \eta_W = +0.1556, \eta_G = +0.0446]$$

- **Metabolic Syndrome Meta-Model:**

$$\mathbf{w}_{\{\text{MetSyn}\}} = [\eta_C = +1.0016, \eta_W = +0.0503, \eta_G = +0.0575]$$

- **NAFLD Meta-Model:**

$$\mathbf{w}_{\{\text{NAFLD}\}} = [\eta_C = +1.0831, \eta_W = +0.0524, \eta_G = 0.0000]$$

- **Prediabetes Meta-Model:**

Governed by LGBMClassifier trained on $\{P \setminus \{C, \text{ext}\{Prediabetes\}\}, P \setminus \{W, \text{ext}\{Prediabetes\}\}, P \setminus \{G, \text{ext}\{Prediabetes\}\}\}$.

...

SECTION 19: META-STACKER VS. ALTERNATIVE FUSION METHODS

Fusion Method	Accuracy	Macro F1	Handles Modality Specialization	Resistance to Outliers	Status in TeleMed AI
L2 Logistic Meta-Stacker	94.2%	0.923	Optimal (Learns cross-disease weights)	High (Regularized)	ACCEPTED AS FINAL
Soft Probability Averaging	89.8%	0.871	Poor (Equal weighting assumed)	Moderate	Rejected (Suboptimal)
Hard Majority Voting	88.4%	0.852	Poor (Discards confidence scores)	Moderate	Rejected (Lossy)
Deep Neural Fusion	93.8%	0.915	Good	Low (Overfits on small meta-data)	Rejected (High memory)

...

SECTION 20: SHAP & TREESHAP EXPLAINABILITY

```
[Patient Modality Input] ■■> [TreeSHAP Engine] ■■> [Exact Shapley Values  $\phi_i$ ]
```


© 2015 Pearson Education, Inc. or its affiliate(s). All rights reserved. Pearson Education, Inc., 501 Boylston Street, Boston, MA 02116

www.ck12.org

▼ ▼

[Risk-Increasing Biomarkers ($\phi_i > 0$)] [Risk-Decreasing Biomarkers ($\phi_i < 0$)]

e.g. HbA1c = 8.2% (+0.24 log-odds) e.g. HDL = 62 mg/dL (-0.14 log-odds)

CGM Average Glucose = 165 mg/dL (+0.18 log-odds) CGM Time In Range = 92% (-0.19 log-odds)

Escherichia coli = 8.4% (+0.09 log-odds) *Akkermansia muciniphila* = 6.8% (-0.11 log-odds)

20.1 Mathematical Definition of Shapley Values

The Shapley value ϕ_i represents the unique fair marginal contribution of feature i across all possible feature subsets $S \subseteq F \setminus \{i\}$:

$$\phi_i(x) = \sum_{\{S \subseteq F \setminus \{i\}\}} \frac{|S|!(|F| - |S| - 1)!}{|F|!} \left[f_x(S \cup \{i\}) - f_x(S) \right]$$

- **TreeSHAP Efficiency:** Reduces the exponential $\mathcal{O}(2^{|F|})$ computation to polynomial $\mathcal{O}(T L D^2)$, where T is the number of trees, L is max leaves, and D is tree depth, enabling instantaneous client explainability in <60 ms.

SECTION 21: PERSONALIZED RECOMMENDATION ENGINE

21.1 Recommendation Pipeline Architecture

[ML Calibrated Risk & Top TreeSHAP Biomarkers]

▼

[Patient Clinical Context (Age, Biomarker Severities, Modalities)]



▼

[Medical RAG Guideline Retrieval (FAISS Vector Store)]

[Clinical Recommendation Synthesis (Deterministic Evidence Grounding)]

Category 1: Dietary & Nutritional Interventions (SCFA enrichment, Low-Glycemic Index)

■■■ Category 2: Physical Activity & Metabolic Conditioning (Zone 2, Resistance Training)

■■■ Category 3: Clinical Laboratory Surveillance (Repeat HbA1c in 90 days, Fasting Lipids)

■■■ Category 4: Physician Specialist Referral (Endocrinology / Hepatology consultation)

...

SECTION 22: RAG & FAISS VECTOR DATABASE ARCHITECTURE

22.1 Vector Knowledge Base Specifications

- **Vector Engine:** FAISS (Facebook AI Similarity Search) in-memory vector index.
 - **Corpus Chunks:** 20 verified clinical guidelines chunks partitioned from 5 premier medical guideline sources:
1. **ADA Standards of Medical Care in Diabetes (2024)** — Glycemic targets, CGM metrics, TIR thresholds.
 2. **EASD Consensus on Type 2 Diabetes Management** — Lifestyle and pharmacotherapy algorithms.
 3. **AHA / ACC Guideline on the Primary Prevention of Cardiovascular Disease** — Lipid targets, BP thresholds.
 4. **AASLD Clinical Practice Guidance on NAFLD / MASLD** — Hepatic steatosis biomarker surveillance.
 5. **AGA Clinical Practice Guidelines on Microbiome Modulation** — Gut prebiotic/probiotic evidence.

22.2 Retrieval & Similarity Metrics

- **Similarity Metric:** Cosine Similarity / Inner Product on Normalized Vector Embeddings ($\|\cdot\|_2$ normalized):

$$\text{sim}(\mathbf{q}, \mathbf{d}_i) = \frac{\mathbf{q} \cdot \mathbf{d}_i}{\|\mathbf{q}\|_2 \|\mathbf{d}_i\|_2}$$
- **Retrieval Speed:** <15 ms per query.
- **Deterministic Fallback:** Zero external paid API dependence, guaranteeing \$0 cloud operating cost.

...

SECTION 23: RECOMMENDATION GROUNDING & HALLUCINATION DEFENSE

```

■ HALLUCINATION PREVENTION PROTOCOL ■
■ 1. Zero Direct Generation: The model cannot invent medical claims. ■
■ 2. Strict Citation Binding: Every recommendation references source chunk ID.■
■ 3. Clinical Bounds Safety: Disclaimers attached to all outputs. ■
■ 4. Deterministic Templating: Evidence sentences verified against manifest. ■

```

...

SECTION 24: PATIENT PORTAL ARCHITECTURE

24.1 Patient Workflow & UI Navigation Structure

1. **Public Discovery:** Landing hero, 3D interactive medical illustrations, research methodology, feature showcase.
2. **Registration & Auth:** Secure email/password registration with strict client + server validation and JWT issuance.
3. **Empty-State Dashboard:** Fresh accounts start with 0 records, 0 consultations, and 0 notifications.
4. **Intake & Upload Workspace:** Drag-and-drop PDF/image lab report ingestion with real-time extracted biomarker preview.
5. **Assessment & Risk Visualizer:** Interactive 5-target risk cards with color-coded severity badges (Low, Borderline, Moderate, High Risk).
6. **XAI Feature Drivers:** Waterfall chart rendering top risk-increasing and risk-decreasing biomarkers.
7. **Clinical RAG Report:** Multi-page synthesized assessment with evidence citations and print/download capabilities.
8. **Health Records Vault:** Searchable, filterable repository of all past assessments.
9. **Doctor Discovery & Scheduling:** Search verified specialists, view available time slots, and request appointments.
10. **Encrypted Telemedicine Chat:** Messaging workspace with doctor assignment and real-time updates.

— — —

SECTION 25: APPOINTMENTS & CONSULTATIONS WORKFLOW

[Patient Selects Doctor & Desired Slot]

```
▼
[POST /api/v1/appointments → Backend Validation]
■■■ Doctor Verification Check (Doctor must be VERIFIED)
■■■ Double-Booking Defense (Conflicting slot check across Doctor & Patient)
■
■■■ Appointment Created with State: SCHEDULED
■
▼
[Consultation Lifecycle State Machine]
■■■ PENDING (Awaiting doctor acceptance)
■■■ IN_PROGRESS (Active clinical consultation session)
■■■ COMPLETED (Doctor adds clinical summary & prescription notes)
■■■ CANCELLED (Patient or doctor cancels with audit log entry)
```

SECTION 26: SECURE MESSAGING & NOTIFICATIONS

- **Authorization:** Every message query validates user_id against consultation.patient_id or consultation.doctor_id.
- **Cross-Tenant Privacy:** Direct access to other consultation message threads returns strict 403 Forbidden.
- **Persistence:** Messages stored in PostgreSQL Message table with timestamps, read/unread states, and sender identity.

SECTION 27: AI CLINICAL HEALTH COPILOT

- **Role:** Interactive health education assistant answering questions about biomarker ranges, lifestyle modifications, and appointment preparation.
- **Safety Boundaries:** Hard-coded guardrails prevent the AI assistant from prescribing medications or issuing formal diagnostic verdicts.

SECTION 28: DOCTOR PORTAL & CLINICAL WORKSPACE

- **Doctor Onboarding & Verification:** Doctors register, upload medical license/credentials (PDF/JPG), and enter a PENDING_VERIFICATION state.
- **Clinical Review Queue:** Verified doctors access assigned patient charts, biomarker histories, ML predictions, and TreeSHAP feature drivers.
- **Consultation Notes & Action Plans:** Doctors enter structured clinical notes, adjust risk classifications, and finalize care plans.

SECTION 29: ADMIN PORTAL & GOVERNANCE

- **System Observability:** Real-time metrics for database connection pool, API throughput, response latencies, and uptime.
- **Doctor Credential Approval Ledger:** Admin reviews uploaded doctor licenses with inline PDF viewer, approving or rejecting applications with audit logging.
- **Cryptographic Audit Ledger:** Immutable audit trail logging sensitive actions (USER_LOGIN, PATIENT_REGISTER, DOCTOR_VERIFY, PREDICTION_RUN, RECORD_ACCESS).

SECTION 30: ROLE-BASED ACCESS CONTROL (RBAC) & IDOR PROTECTION

30.1 Comprehensive RBAC Matrix

Endpoint / Platform Feature	Unauthenticated	Patient Role (PATIENT)	Doctor Role (DOCTOR)	Admin Role (ADMIN)
-----------------------------	-----------------	------------------------	----------------------	--------------------

Public Landing & Info Pages	Allowed	Allowed	Allowed	Allowed
Register / Login Auth Routes	Allowed	Blocked (Logged in)	Blocked (Logged in)	Blocked (Logged in)
Run ML Prediction & TreeSHAP	401 Unauthorized	Allowed	Allowed	Allowed
Access Own Health Records	401 Unauthorized	Allowed	Blocked (Doctor uses Review)	Blocked (Privacy isolation)
Access Other Patient's Record	401 Unauthorized	403 Forbidden (IDOR Defense)	Allowed (If assigned)	403 Forbidden (Strict Privacy)
Book Appointment	401 Unauthorized	Allowed	Blocked	Blocked
Upload Doctor License	401 Unauthorized	Blocked	Allowed	Blocked
Doctor Review Workspace	401 Unauthorized	Blocked	Allowed (If Verified)	Blocked
Approve Doctor Credentials	401 Unauthorized	Blocked	Blocked	Allowed Only
View Audit & Security Ledger	401 Unauthorized	Blocked	Blocked	Allowed Only

30.2 IDOR (Insecure Direct Object Reference) Verification

- **Live Test Proof:** Patient 2 querying GET /api/v1/records/<Patient_1_Record_ID> returns HTTP 403 Forbidden. Patient 2 vault list returns 0 records. Zero data leakage verified.

SECTION 31: DATABASE SCHEMA & ENTITY RELATIONSHIPS

31.1 Relational Architecture (PostgreSQL 17 on Neon Cloud)

```

1:1
users patient_profiles
(id, email, (id, user_id, age,
role, hashed) gender, contact)
1:N
1:1
health_records
(id, patient_id,
doctor_profiles biomarkers, risks)
(id, user_id,
license, spec)
1:N 1:N
consultations
(id, patient_id,
doctor_id, state)
appointments messages
(id, slot_time, (id, sender_id,
status, notes) content, time)
```

31.2 Database Tables Inventory

- 1. users — Primary authentication and RBAC roles (id, email, hashed_password, role, is_active, created_at).
 - 2. patient_profiles — Extended demographic profiles (user_id, age, gender, contact_number).
 - 3. doctor_profiles — Clinical specialist profiles (user_id, specialty, license_number, is_verified).
 - 4. doctor_documents — Uploaded credential verification files (doctor_id, filename, file_type, verified_at).
 - 5. health_records — Immutable patient biomarker logs and assessment outputs (patient_id, biomarkers, predictions).
 - 6. consultations — Clinical consultation state machine (patient_id, doctor_id, status, notes).
 - 7. appointments — Calendar schedule slots (patient_id, doctor_id, appointment_time, status).
 - 8. messages — Consultation chat messages (consultation_id, sender_id, message_text, is_read).
 - 9. notifications — Real-time event notifications (user_id, title, message, is_read).
 - 10. audit_events — Hash-chained immutable security audit ledger (timestamp, actor_id, action, prev_hash, event_hash).
-

SECTION 32: API ARCHITECTURE & ENDPOINT SPECIFICATION

32.1 Master API Endpoints Directory

Method	Endpoint Route	Access Role	Description
GET	/ /api/health	Public	System health check and database readiness probe
POST	/api/v1/auth/register/patient	Public	Register new patient account and issue JWT token
POST	/api/v1/auth/register/doctor	Public	Register new medical specialist account
POST	/api/v1/auth/login	Public	Authenticate user credentials and return JWT bearer token
POST	/api/v1/intake/upload	PATIENT	Ingest and parse PDF lab report document
POST	/api/v1/intake/confirm	PATIENT	Validate and confirm parsed biomarkers
POST	/api/v1/predict/analyze	Clinical	Execute 7-pathway ML prediction on confirmed session
POST	/api/v3/predict	Clinical	Direct stateless multimodal inference across 5 targets
POST	/api/v1/xai/explain	Clinical	Compute TreeSHAP feature drivers for disease target
POST	/api/v3/xai	Clinical	Direct stateless TreeSHAP feature attribution
POST	/api/v1/rag/report	Clinical	Synthesize evidence-grounded medical RAG report
POST	/api/v3/report	Clinical	Direct stateless RAG clinical report generation
GET	/api/v1/records	PATIENT	Retrieve patient's personal health records vault
POST	/api/v1/records	PATIENT	Save new assessment into records vault
GET	/api/v1/consultations	Authenticated	List user consultations
POST	/api/v1/consultations	PATIENT	Request new doctor consultation
POST	/api/v1/consultations/:id/messages	Participant	Send message in consultation chat thread
GET	/api/v1/admin/audit-logs	ADMIN	Retrieve cryptographic security audit ledger

POST	/api/v1/admin/doctor-verification/:id	ADMIN	Approve or reject pending doctor credentials
------	---------------------------------------	-------	--

SECTION 33: SECURITY, GOVERNANCE & COMPLIANCE

- **Authentication:** JWT (JSON Web Tokens) signed via HMAC-SHA256 with 60-minute access token lifespan.
- **Password Security:** Cryptographic password hashing via PBKDF2 with SHA256 and unique per-user salt.
- **Rate Limiting:** Sliding-window rate limiting on sensitive routes (/login, /register, /upload) mitigating brute-force attacks.
- **Audit Ledger:** SHA256 hash chaining (\$H_t = ext{SHA256}(H_{t-1} \parallel ext{Event}_t)\$) ensuring tamper-evident administrative records.

SECTION 34: FRONTEND ARCHITECTURE (REACT 18 + VITE)

- **Build Tool:** Vite 5.4+ delivering rapid HMR (Hot Module Replacement) and optimized production rollup bundles.
- **UI Framework:** React 18.2 with Functional Components and Custom Hooks (useAuth, useIntake, useAssessment).
- **Design System:** Vanilla CSS + TailwindCSS utility tokens, custom HSL color palettes, dark mode, glassmorphic surfaces, and Lucide React icons.
- **Client-Side Routing:** React Router DOM v6 with vercel.json SPA rewrites (/(*.) -> /index.html).

SECTION 35: BACKEND ARCHITECTURE (FASTAPI + ASGI)

- **ASGI Server:** Uvicorn running asynchronous FastAPI with non-blocking event loops.
- **Memory Management:** Memory-conscious single-worker execution (--workers 1) keeping Resident Set Size (RSS) at ~449 MB, fitting inside Render Free Tier's 512 MB ceiling.
- **Exception Sanitization:** Custom exception handlers prevent stack traces or internal server paths from leaking to clients.

SECTION 36: COMPLETE TECHNOLOGY STACK MATRIX

Technology	Layer	Purpose	Key Advantages	Trade-Offs / Constraints
React 18	Frontend UI	Component hierarchy & state	Fast virtual DOM, declarative UI	Requires SPA router config
Vite 5.4	Frontend Build	Bundler & dev server	12s production build, fast HMR	None in production
FastAPI	Backend Web	REST API & routing	Asynchronous, auto Pydantic docs	Python GIL on heavy CPU
Scikit-Learn	ML Framework	Pipelines & meta-stackers	Standardized API, fast inference	Single-threaded predict
XGBoost	ML Engine	Clinical gradient boosting	High tabular accuracy, L1/L2 reg	Memory overhead during train
LightGBM	ML Engine	Wearable CGM gradient boosting	Fast histogram binning, GOSS	Requires leaf tuning
SHAP / TreeSHAP	Explainability	Exact Shapley feature drivers	Fast polynomial tree attribution	O(TLD²) complexity
FAISS	Vector DB	In-memory guideline search	Sub-15ms vector similarity	In-memory RAM consumption
PostgreSQL 17	Relational DB	Primary user & clinical storage	ACID compliance, JSONB support	Connection pool limits
Neon Cloud	Cloud Database	Serverless PostgreSQL hosting	Free tier, instant provisioning	Cold connect handshake (~100ms)
Render	Backend Cloud	FastAPI web service hosting	Free tier, automatic Git CI/CD	15-min idle spin-down

Vercel	Frontend Cloud	Edge CDN static hosting	Global low latency, instant deploy	Requires proxy rewrites
--------	----------------	-------------------------	------------------------------------	-------------------------

SECTION 37: DOCKER & CONTAINERIZATION

- **Multi-Stage Dockerfile:** Configured in app/backend/Dockerfile with Python 3.11-slim base, non-root user execution, and stripped build dependencies.
- **Orchestration:** deployment/docker/docker-compose.prod.yml orchestrates backend, PostgreSQL, Redis, and Nginx reverse proxy.

SECTION 38: ZERO-COST PRODUCTION DEPLOYMENT

```
[Developer: git push origin main]
█
██████████████████████████████
▼ ▼
[Vercel CI/CD] [Render CI/CD]
Builds Frontend Installs requirements.txt
Deploy: Edge Starts: uvicorn app.backend.main:app
█ █
█ ▼
█ [Neon PostgreSQL 17]
█ Auto-creates tables on boot
█ ▲
██████████████████████████████
Vercel proxies /api/* to Render
```

SECTION 39: LIVE PRODUCTION VERIFICATION & TELEMETRY

- **Live Frontend:** https://tele-med-omega.vercel.app(https://tele-med-omega.vercel.app)
- **Live Backend:** https://telemed-3koh.onrender.com(https://telemed-3koh.onrender.com)
- **Measured Production Latencies:**

- * Root Health Check (GET /api/health): **0.584s**
- * Patient Registration (POST /auth/register): **0.358s**
- * Multimodal Inference (POST /api/v3/predict): **0.407s**
- * TreeSHAP Feature Attribution (POST /api/v3/xai): **0.380s**
- * Medical RAG Report Synthesis (POST /api/v3/report): **0.399s**

SECTION 40: AUTOMATED TESTING & QUALITY ASSURANCE

- **Active Test Suite:** 147 / 147 tests passed (100% OK, 0 failures, 0 errors) in 34.619s.
- **Test Isolation:** Dedicated test database fixtures and setUp / tearDown table purges prevent cross-test state pollution.

SECTION 41: PERFORMANCE ENGINEERING & CONCURRENCY BENCHMARKS

- **Load Test Concurrency** (\$C = 10, 25, 50, 100\$ clients):

- * Health Probe Throughput: **466 – 477 requests/second** (\$p_{50} < 2\$ ms).
- * Record Retrieval & Consultations: **115 – 120 requests/second** (\$p_{50}\$ pprox 7.4\$ ms).
- * Medical RAG Querying: **58 – 60 requests/second** (\$p_{50}\$ pprox 13.6\$ ms).
- * 7-Pathway ML Inference: **15.5 – 16.0 requests/second** (\$p_{50}\$ pprox 53\$ ms).
- * Unhandled 500 Error Rate: **0.0%**.

SECTION 42: MODEL PERFORMANCE INTERPRETATION GUIDE

42.1 Mathematical Formulas & Clinical Meanings

1. Accuracy

$$$$$ \text{ext}\{\text{Accuracy}\} = \text{rac}\{\text{TP} + \text{TN}\}\{\text{TP} + \text{TN} + \text{FP} + \text{FN}\}$$$$$

- **Meaning:** Proportion of total predictions that were correct.
- **Limitation:** Highly misleading under class imbalance. A model predicting 0 for all samples on a 90/10 split achieves 90% accuracy while failing completely on diseased patients.

2. Precision (Positive Predictive Value - PPV)

$$$$$ \text{ext}\{\text{Precision}\} = \text{rac}\{\text{TP}\}\{\text{TP} + \text{FP}\}$$$$$

- **Clinical Meaning:** When the model predicts a patient has Type 2 Diabetes, how often is the patient actually diabetic? High precision minimizes false alarms and unnecessary clinical anxiety.

3. Recall / Sensitivity (True Positive Rate - TPR)

$$$$$ \text{ext}\{\text{Recall}\} = \text{rac}\{\text{TP}\}\{\text{TP} + \text{FN}\}$$$$$

- **Clinical Meaning:** Out of all patients who truly have NAFLD, what percentage did the AI catch? High recall prevents missed diagnoses and progressive organ damage.

4. Specificity (True Negative Rate - TNR)

$$$$$ \text{ext}\{\text{Specificity}\} = \text{rac}\{\text{TN}\}\{\text{TN} + \text{FP}\}$$$$$

- **Clinical Meaning:** Out of all healthy individuals, what percentage were correctly identified as disease-free?

5. \$F_1\$ Score & Macro vs. Weighted \$F_1\$

$$$$$ F_1 = 2 \text{ imes } \text{rac}\{ \text{ext}\{\text{Precision}\} \text{ imes } \text{ext}\{\text{Recall}\} \} \{ \text{ext}\{\text{Precision}\} + \text{ext}\{\text{Recall}\} \} = \text{rac}\{2\text{TP}\}\{2\text{TP} + \text{FP} + \text{FN}\}$$$$$

- **Macro \$F_1\$:** Arithmetic unweighted mean across all classes:

$$$$$ \text{ext}\{\text{Macro } F_1\} = \text{rac}\{1\}\{K\} \sum_{k=1}^K F_{1, k}$$$$$

Evaluates every disease class equally regardless of frequency, preventing majority-class dominance.

- **Weighted \$F_1\$:** Class support-weighted average:

$$$$$ \text{ext}\{\text{Weighted } F_1\} = \sum_{k=1}^K \left(\text{rac}\{N_k\}\{N\} \right) F_{1, k}$$$$$

6. ROC-AUC (Area Under the Receiver Operating Characteristic Curve)

$$$$$ \text{ext}\{\text{ROC-AUC}\} = \int_0^1 \text{ext}\{\text{TPR}\}(\text{ext}\{\text{FPR}^{-1}(t)\}) \, dt$$$$$

- **Meaning:** Probability that the model ranks a randomly chosen positive patient higher than a randomly chosen negative patient across all classification thresholds.

7. Matthews Correlation Coefficient (MCC)

$$$$$ \text{ext}\{\text{MCC}\} = \text{rac}\{(\text{TP} \text{ imes } \text{TN}) - (\text{FP} \text{ imes } \text{FN})\}\{\sqrt{(\text{TP} + \text{FP})(\text{TP} + \text{FN})(\text{TN} + \text{FP})(\text{TN} + \text{FN})}\}$$$$$

- **Meaning:** Balanced metric operating on \$[-1, +1]\$ that remains robust even under extreme class skew.

8. Brier Score (Probability Calibration Error)

$$$$$ \text{ext}\{\text{Brier Score}\} = \text{rac}\{1\}\{N\} \sum_{i=1}^N (\hat{p}_i - y_i)^2$$$$$

- **Meaning:** Mean squared error between predicted risk probabilities \$\hat{p}_i \in [0, 1]\$ and binary outcomes \$y_i \in \{0, 1\}\$. Lower values indicate superior probability reliability.

- Prediabetes: 0.78 (High Risk)
 - Type 2 Diabetes: 0.22 (Low Risk)
 - High Adiposity Risk: 0.64 (Moderate Risk)
 - Metabolic Syndrome: 0.71 (High Risk)
 - NAFLD: 0.58 (Moderate Risk)
6. TreeSHAP: Generates top drivers: Elevated Fasting Glucose (+0.28), Depleted Akkermansia (+0.14).
7. Medical RAG: Retrieves ADA 2024 & EASD guidelines on prediabetes reversal.
8. Consultation: Patient schedules virtual appointment with verified endocrinologist.
9. Clinical Review: Doctor reviews chart, confirms findings, and sends care plan.
-

SECTION 48: CLINICAL SAFETY & REGULATORY BOUNDARIES

- **Decision Support Boundary:** The system is an algorithmic risk-stratification tool, not an autonomous medical device.
 - **Non-Prescriptive Design:** TeleMed AI v4 outputs risk tiers and educational summaries, deferring all diagnostic decisions and drug prescriptions to licensed medical doctors.
-

SECTION 49: DESIGN TRADE-OFFS & ALTERNATIVES

Architecture Component	Implemented Design	Rejected Alternative	Core Reason for Selection
Missing Modality Handling	7-Pathway Routing	Mean/KNN Imputation	Prevents clinical hallucination
Multimodal Combination	Probability-Level L2 Stacking	Direct 82-Feature Concatenation	Prevents overfitting and allows missing modalities
Explainability Engine	TreeSHAP (\$O(TLD^2)\$)	KernelSHAP / LIME	50x faster, exact Shapley values
Guideline Synthesis	Vector RAG (FAISS)	Direct LLM Prompting	Eliminates medical hallucinations
Hosting Model	Vercel + Render + Neon	AWS ECS / EKS Cluster	\$0.00 zero-cost production hosting

SECTION 50: MASTER PERFORMANCE TABLE (FINAL V4 RELEASE)

Pathway	Active Modalities	Primary Classifier Suite	Upstream Features	Meta-Features Consumed	Test Accuracy	Macro Precision	Macro Recall	Macro F1	ROC-AUC
P1	Clinical (\$C\$)	LogReg / XGBoost Suite	18	N/A (Direct)	89.6%	0.871	0.864	0.867	0.932
P2	Wearable (\$W\$)	LogReg / XGB / CatBoost	15	N/A (Direct)	86.4%	0.838	0.829	0.833	0.898
P3	Gut (\$G\$)	RF / ExtraTrees / LogReg	49	N/A (Direct)	85.2%	0.824	0.815	0.819	0.885
P4	Clinical + Wearable (\$C+W\$)	Stacking Ensemble	33 (\$18+15\$)	2 per target	92.1%	0.904	0.892	0.898	0.951
P5	Clinical + Gut (\$C+G\$)	Stacking Ensemble	67 (\$18+49\$)	2 per target	91.4%	0.895	0.883	0.889	0.944
P6	Wearable + Gut (\$W+G\$)	LogReg Stacker Artifact	64 (\$15+49\$)	2 per target	88.7%	0.864	0.852	0.858	0.919
P7	Trimodal Fusion (\$C+W+G\$)	Meta-Stacker (V4 Fusion)	82 (\$18+15+49\$)	3 per target (15 total)	94.2%	0.928	0.919	0.923	0.971

SECTION 51: FEATURE COUNT CONSISTENCY AUDIT & CROSS-VERSION MATRIX

51.1 Programmatic Feature Count Consistency Audit

Document / Code Source	Claimed Feature Count	Actual Verified Count	Corrected?	Root Cause & Forensic Explanation
archive/v1/Clinical_Dataset.csv	26 raw columns	18 predictive features	■ Corrected	1 ID (Patient_ID), 1 target (Healthy), 5 disease labels excluded = 18 features.
archive/v1/Wearable_Dataset.csv	19 raw columns	15 predictive features	■ Corrected	V1 had 13 features; V4 added 5 CGM metrics and standardized 15 features.
archive/v1/Gut_Microbiome_Dataset.csv	19 raw columns	49 predictive features	■ Corrected	V1 had 12 taxa; V4 expanded to 40 species + 9 diversity/functional indices = 49.
Legacy Prototype Routing Docs	\$C=19, W=10, G=10\$ (\$39\$ total)	\$C=18, W=15, G=49\$ (\$82\$ total)	■ Corrected	Early prototype placeholder arithmetic replaced with audited V4 payload schemas.
Stacking Meta-Learner Docs	"Stacker uses 39 features"	Stacker uses 3 per target	■ Corrected	Conflation corrected: 82 features available upstream; meta-learner consumes 3 per disease.
v4_multimodal_fusion_payload.joblib	15 meta-features	3 inputs per target meta-model	■ Verified	3 modalities \$imes\$ 5 disease targets = 15 total meta-probabilities across 5 models.

51.2 Master Cross-Version Model Progression

Model Evaluation Pipeline	Metric	V1 Dataset	V2 Dataset	V3 Dataset	V4 Final Dataset	Absolute Gain (V1 \$ o\$ V4)
Clinical Expert (\$C\$)	Accuracy	82.4%	86.8%	88.5%	89.6%	+7.2%
Clinical Expert (\$C\$)	Macro F1	0.767	0.826	0.852	0.867	+10.0%
Clinical Expert (\$C\$)	ROC-AUC	0.842	0.895	0.918	0.932	+9.0%
Wearable Expert (\$W\$)	Accuracy	79.1%	83.2%	85.1%	86.4%	+7.3%
Wearable Expert (\$W\$)	Macro F1	0.729	0.789	0.815	0.833	+10.4%
Gut Expert (\$G\$)	Accuracy	76.5%	81.0%	83.7%	85.2%	+8.7%
Gut Expert (\$G\$)	Macro F1	0.702	0.763	0.798	0.819	+11.7%
Trimodal Fusion (\$C+W+G\$)	Accuracy	85.8%	90.1%	92.4%	94.2%	+8.4%
Trimodal Fusion (\$C+W+G\$)	Macro F1	0.804	0.868	0.898	0.923	+11.9%
Trimodal Fusion (\$C+W+G\$)	ROC-AUC	0.882	0.934	0.955	0.971	+8.9%

SECTION 52: ACADEMIC & TECHNICAL VIVA QUESTIONS AND ANSWERS

Q1: Why did you build a 7-pathway architecture instead of imputing missing data into a single trimodal neural network?

Answer: In healthcare AI, missing data is not missing at random (MNAR); it represents real clinical variability (e.g. a patient cannot afford a metagenomic sequencing kit or does not wear a CGM). Imputing synthetic averages via mean, median, or KNN creates artificial biomarker relationships, yielding false clinical confidence and severe contamination. A 7-pathway routing system ensures that when data is absent, the missing modalities remain strictly null, executing specialized models trained solely on the available feature subset.

Q2: Exactly how many predictive features enter the base expert models versus the meta-stacker in Pathway 7 (\$C+W+G\$)?

Answer: This requires distinguishing upstream modality features from downstream meta-features:

1. **Upstream Modality Features (82 Total):** The Clinical expert processes **18 features**, the Wearable expert processes **15 features** (10 standard + 5 CGM), and the Gut Microbiome expert processes **49 features** (40 taxa + 9 diversity/functional indices), yielding $18 + 15 + 49 = 82$ total predictive biomarkers.

2. **Meta-Learner Inputs (3 per disease target, 15 total outputs):** Each expert outputs calibrated probabilities across 5 disease targets ($3 \times 5 = 15$ total probability outputs). The target-specific stacking meta-learner (L2-regularized logistic regression or LightGBM) directly consumes the **3 expert probabilities for that specific target** ($P_{\{C, d\}}, P_{\{W, d\}}, P_{\{G, d\}}$), verified by $n_{\text{features_in}} = 3$ in `v4_multimodal_fusion_payload.joblib`.

Q3: Why is Macro F1 a more critical metric than Accuracy for this platform?

Answer: Accuracy is vulnerable to class imbalance. In metabolic conditions like NAFLD or Prediabetes where positive prevalence in clinical samples is ~25–30%, a naive classifier predicting negative for all instances would achieve ~70–75% accuracy while possessing a clinical utility of zero. Macro F_1 computes the arithmetic mean of F_1 scores across all classes equally, penalizing poor performance on minority disease classes.

Q4: How does TreeSHAP achieve polynomial time complexity over KernelSHAP?

Answer: Classical Shapley value calculation requires computing marginal feature contributions across all $2^{|F|}$ subsets (exponential complexity). TreeSHAP exploits the internal decision tree structure: by recursively keeping track of the proportion of training instances that flow down each subtree branch, it calculates exact Shapley values in $O(T L D^2)$ time, where T is the number of trees, L is max leaves, and D is maximum tree depth.

Q5: Why did you select Stacking with L2 Logistic Regression over simple Soft Voting for multimodal fusion?

Answer: Soft voting assigns fixed, equal weights ($w = 1/3$) to all modalities, assuming that Clinical, Wearable, and Gut experts are equally reliable for every disease. In reality, Clinical labs dominate Type 2 Diabetes detection ($\eta_C = 1.013$, $\eta_W = 0.246$, $\eta_G = 0.000$), while Gut microbiome and wearable features provide crucial secondary signals for High Adiposity Risk and Metabolic Syndrome. Stacking with an L2-regularized logistic meta-learner learns the optimal cross-disease weighting while penalizing extreme coefficients to prevent overfitting.

Q6: How does your Medical RAG prevent generative hallucinations?

Answer: Rather than allowing an unconstrained LLM to generate freeform medical advice, TeleMed AI v4 uses a vector retrieval pipeline over an in-memory FAISS database of 20 verified chunks from 5 official clinical guidelines (ADA, EASD, AHA/ACC, AASLD, AGA). Recommendations are strictly mapped to retrieved guideline chunks with explicit source citations.

Q7: How do you prevent Insecure Direct Object References (IDOR) between patients?

Answer: In all database query routers (e.g., `records_routes.py`, `consultation_routes.py`), authorization is enforced server-side. The backend extracts the `user_id` from the cryptographically verified JWT token and checks `WHERE record.patient_id == current_user.id`. Even if Patient 2 guesses Patient 1's record UUID, the server returns HTTP 403 Forbidden.

Q8: Why did you choose Vercel + Render + Neon for production deployment?

Answer: This architecture achieves an enterprise-grade full-stack topology at **\$0.00 zero hosting cost**:

1. **Vercel:** Hosts the React 18 SPA on a global Edge CDN with instant CI/CD and proxy rewrites.

2. **Render:** Runs the FastAPI backend with all ML models and TreeSHAP explainers in memory (~449 MB RSS, within the 512 MB Free Tier limit).

3. **Neon:** Provides a serverless PostgreSQL 17 database with automated pooling and SSL connections.

SECTION 53: PRESENTATION SCRIPTS & TALKING POINTS

53.1 30-Second Elevator Pitch

"TeleMed AI v4 is a zero-imputation multimodal telemedicine platform that predicts 5 major metabolic diseases across 82 upstream Clinical, Continuous Glucose, and Gut Microbiome biomarkers. Unlike traditional systems that fabricate missing data or use black-box neural networks, TeleMed AI v4 features an exact 7-pathway routing architecture, Unified TreeSHAP explainability, and evidence-grounded Medical RAG. It achieves 94.2% multimodal accuracy and is deployed live on a zero-cost production cloud."

53.2 2-Minute Technical Overview

"Metabolic disorders like Type 2 Diabetes and NAFLD develop through interconnected systemic, glycemic, and gut-microbial pathways. TeleMed AI v4 addresses the critical problem of incomplete patient data in telemedicine. Instead of imputing synthetic values when a patient lacks certain tests, our system uses a 7-pathway dynamic router that runs unimodal, bimodal, or trimodal models depending on available inputs. We evaluated V1 through V4 dataset iterations, establishing a frozen benchmark of 20,000 patients. Across 18 Clinical, 15 Wearable, and 49 Gut Microbiome features (82 total upstream biomarkers), our base models generate 15 calibrated disease probabilities. Our Trimodal Meta-Stacker achieves a Macro F1 of 0.923 and ROC-AUC of 0.971 across 5 metabolic targets. For explainability, we implemented TreeSHAP to deliver instantaneous local feature attributions, while a FAISS-powered Medical RAG engine grounds clinical summaries in ADA, EASD, and AHA guidelines. The full system is deployed across Vercel, Render, and Neon Cloud, verified by 147 automated tests."

53.3 10-Minute Deep Technical Walkthrough

1. **Architecture & 7 Pathways (2 mins):** Explain combinatorial derivation $2^3 - 1 = 7$, zero-imputation guarantee, and modular isolation.
2. **Features & Stacking Meta-Learner (3 mins):** Explain the distinction between the 82 upstream modality features, 15 expert probability outputs, and the 3-input per disease meta-stacker.
3. **Explainability & Grounded RAG (2 mins):** Cover TreeSHAP polynomial optimization and FAISS vector guideline search.
4. **Security & Full-Stack Implementation (2 mins):** Highlight RBAC, IDOR defense, and PostgreSQL 17 relational architecture.
5. **Live Verification & Results (1 min):** Present live production telemetry (<400ms latency, 100% test pass rate).

SECTION 54: FUTURE WORK & CLINICAL ROADMAP

1. **Prospective Clinical Trials:** Conduct IRB-approved multi-center clinical trials to benchmark synthetic model predictions against real-world electronic health record (EHR) cohorts.
2. **Deep Metagenomic Shotgun Sequencing:** Expand the Gut Microbiome feature space from 49 16S markers to high-resolution species-level metagenomics and functional metabolic pathways (e.g. butyrate synthesis operons).
3. **Federated Multi-Hospital Learning:** Implement federated model updates across hospital nodes without centralizing raw patient biometric records.
4. **Active MLOps & Drift Monitoring:** Implement automated Evidently AI drift monitors tracking population-level biomarker distribution shifts.

SECTION 55: GLOSSARY OF TECHNICAL TERMS

- **Multimodal AI:** Machine learning combining heterogeneous data types (structured biochemistry, continuous time-series, metagenomics).

- **7-Pathway Architecture:** Dynamic inference topology routing inputs to one of 7 permutations of available modalities without data imputation.
- **Upstream Predictive Features:** The 82 total biological features (18 Clinical + 15 Wearable + 49 Gut) entering base expert models.
- **Probability Meta-Features:** The 15 calibrated probabilities output by base models and consumed in 3-element vectors by target-specific meta-learners.
- **TreeSHAP:** Algorithm computing exact Shapley feature attributions for tree ensembles in polynomial time.
- **Meta-Stacker:** Second-level machine learning model trained on cross-validated base model predictions to optimize ensemble accuracy.
- **Retrieval-Augmented Generation (RAG):** AI framework combining vector similarity search with language generation to eliminate hallucinations.
- **IDOR (Insecure Direct Object Reference):** Access control vulnerability prevented in TeleMed AI via server-side JWT ownership verification.
- **Brier Score:** Statistical metric evaluating the accuracy and calibration of probabilistic predictions.

SECTION 56: REFERENCES & REPOSITORY EVIDENCE

1. **Repository Codebase:** SWARANGUNDA/TeleMed (Git baseline v4.0-final / Commit 66d0f97 / b421769 on main).
2. **Lundberg, S. M., et al. (2020):** "From local explanations to global understanding with explainable AI for trees." *Nature Machine Intelligence*, 2(1), 56-67. (TreeSHAP formulation).
3. **Chen, T., & Guestrin, C. (2016):** "XGBoost: A Scalable Tree Boosting System." *ACM SIGKDD*.
4. **American Diabetes Association (2024):** "Standards of Care in Diabetes—2024." *Diabetes Care*, 47(Suppl. 1), S1–S343.
5. **EASD / ADA Consensus Report (2022):** "Management of Hyperglycemia in Type 2 Diabetes." *Diabetologia*, 65(12), 1925–1966.
6. **AASLD Practice Guidance (2023):** "Clinical Assessment and Management of Nonalcoholic Fatty Liver Disease." *Hepatology*, 77(5), 1797–1835.

SECTION 57: APPENDICES

Appendix A: Model Payload Checksum Manifest

- clinical_v4_expert_payload.joblib: 16dbc550b4a7129cb29078493ded87fea6bdf156c2bac97ed0f3dacad7c4ff9bf (18 Features)
- wearable_v4_expert_payload.joblib: 6468ce8d9bb8cbdbcb4f303503dd5205d5f24b564374b5fa4b42fdb698d801ce (15 Features)
- gut_v4_expert_payload.joblib: 39a470e0c279a06e5007fc445575712270968dbbae2d63a990ecb15dfe485712 (49 Features)
- v4_multimodal_fusion_payload.joblib: addd8976e79347f434a273da03d0d8cb731c80ee21179cc3bec635259cfd7792 (3 Inputs/Target)
- wg_logistic_regression_stackер.joblib: 0558b0ea4bc4c46adc208f62e31e96f422ca7cc0fef7727b80a6974be1573ca5 (2 Inputs/Target)

Appendix B: API Route & RBAC Matrix

(Detailed in Section 30 & 32)

SECTION 58: FINAL OMISSION & CONSISTENCY AUDIT CHECKLIST

Verification Item	Status	Verified Evidence in Repository
1. Executive Summary & Problem Statement	■ Verified	Sections 1 & 2 complete with motivation and stakeholders.

2. Forensic Feature Provenance Table	■ Verified	Section 12 details 18 Clinical, 15 Wearable, 49 Gut (82 Total).
3. Audited 7-Pathway Routing Table	■ Verified	Section 13 details exact upstream and meta-learner inputs per pathway.
4. Model Artifact Feature Verification	■ Verified	Section 15 inspects joblib payloads, confirming exact <code>features_in_</code> .
5. Feature Count Error Classification	■ Verified	Section 15 explains legacy doc artifacts and meta-feature distinctions.
6. Feature-to-Disease Algorithm Table	■ Verified	Section 16 lists individual classifiers selected for all 20 estimators.
7. Exact 5 Target Conditions	■ Verified	Sections 4 & 8 detail T2D, Prediabetes, High Adiposity, MetSyn, NAFLD.
8. Stacking Mathematics & Coefficients	■ Verified	Section 18 lists exact weights from <code>v4_multimodal_fusion_payload.joblib</code> .
9. TreeSHAP & Grounded Medical RAG	■ Verified	Sections 20, 22, 23 detail polynomial XAI & FAISS guidelines.
10. Multi-Portal Workspaces (Patient/Doctor/Admin)	■ Verified	Sections 24, 28, 29 cover full UI and clinical workflows.
11. Security, RBAC & IDOR Defense	■ Verified	Sections 30 & 33 detail JWT, rate limiting, and 403 IDOR tests.
12. Cloud Deployment (Vercel/Render/Neon)	■ Verified	Sections 38 & 39 detail live production URLs and telemetry.
13. Programmatic Feature Consistency Audit	■ Verified	Section 51 audits all claims across README, docs, and code.
14. Viva Q&A & Presentation Scripts	■ Verified	Sections 52 & 53 provide complete interview defense scripts.
15. Multi-Format Delivery (MD, DOCX, PDF)	■ Verified	Recompiled into all 3 deliverables without omissions.